

Chapter 5

The LC-3

Instruction Set Architecture

ISA = All of the *programmer-visible* components and operations of the computer

- **memory organization**
 - address space -- how many locations can be addressed?
 - addressability -- how many bits per location?
- **register set**
 - how many? what size? how are they used?
- **instruction set**
 - opcodes
 - data types
 - addressing modes

ISA provides all information needed for someone that wants to write a program in **machine language** (or translate from a high-level language to machine language).

LC-3 Overview: Memory and Registers

Memory

- address space: **2^{16}** locations (16-bit addresses)
- addressability: **16 bits**

Registers

- temporary storage, accessed in a single machine cycle
 - accessing memory generally takes longer than a single cycle
- eight general-purpose registers: **R0 - R7**
 - each **16 bits wide**
 - how many bits to uniquely identify a register?
- other registers
 - not directly addressible, but used by (and affected by) instructions
 - **PC** (program counter), **condition codes**

LC-3 Overview: Instruction Set

Opcodes

- 15 opcodes
- *Operate* instructions: ADD, AND, NOT
- *Data movement* instructions: LD, LDI, LDR, LEA, ST, STR, STI
- *Control* instructions: BR, JSR/JSRR, JMP, RTI, TRAP
- some opcodes set/clear *condition codes*, based on result:
 - N = negative, Z = zero, P = positive (> 0)

Data Types

- 16-bit 2's complement integer

Addressing Modes

- How is the location of an operand specified?
- non-memory addresses: *immediate*, *register*
- memory addresses: *PC-relative*, *indirect*, *base+offset*

Operate Instructions

Only three operations: **ADD, AND, NOT**

Source and destination operands are **registers**

- These instructions do not reference memory.
- ADD and AND can use “immediate” mode, where one operand is hard-wired into the instruction.


Will show **dataflow diagram** with each instruction.

- illustrates when and where data moves to accomplish the desired operation

NOT (Register)

	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
NOT	1	0	0	1	Dst			Src			1	1	1	1	1	1




NOT Dst, Src



*Note: Src and Dst
could be the same register.*



ADD/AND (Register)

this zero means "register mode"



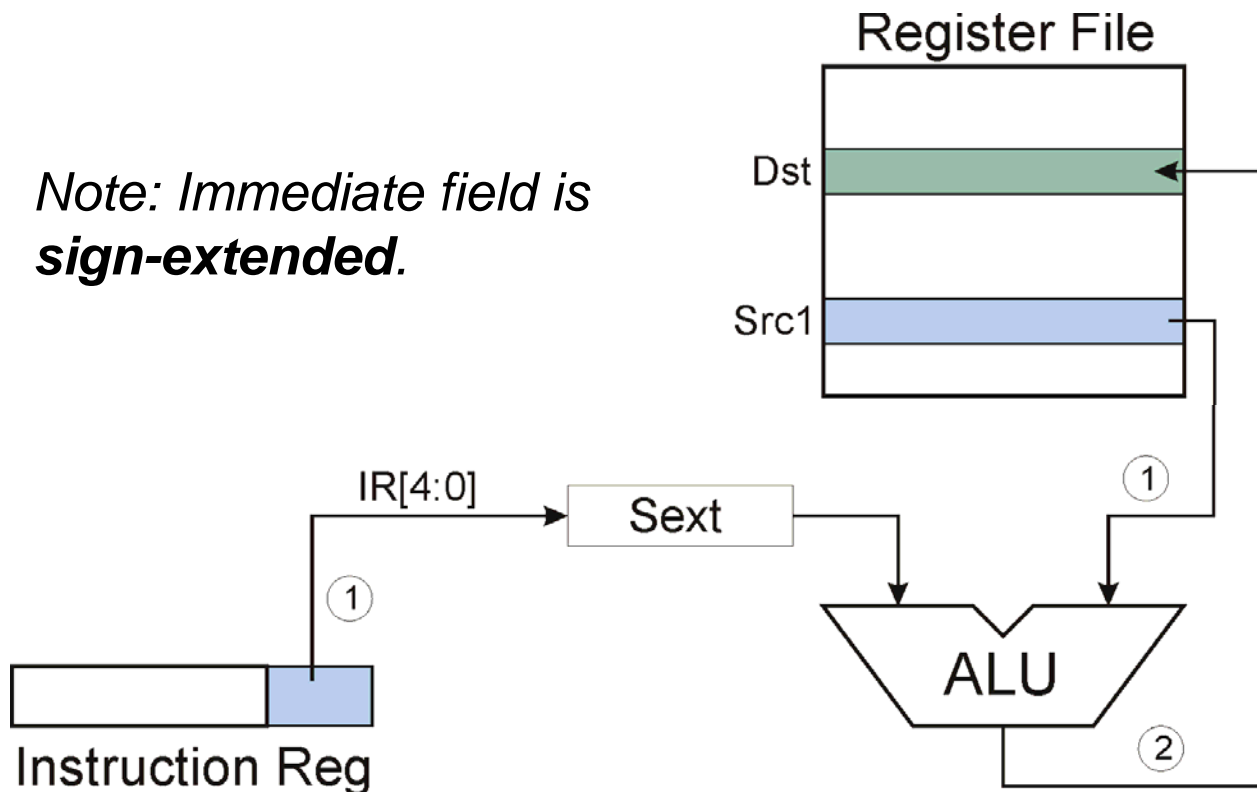
ADD/AND (Immediate)

this one means "immediate mode"



*Note: Immediate field is **sign-extended**.*

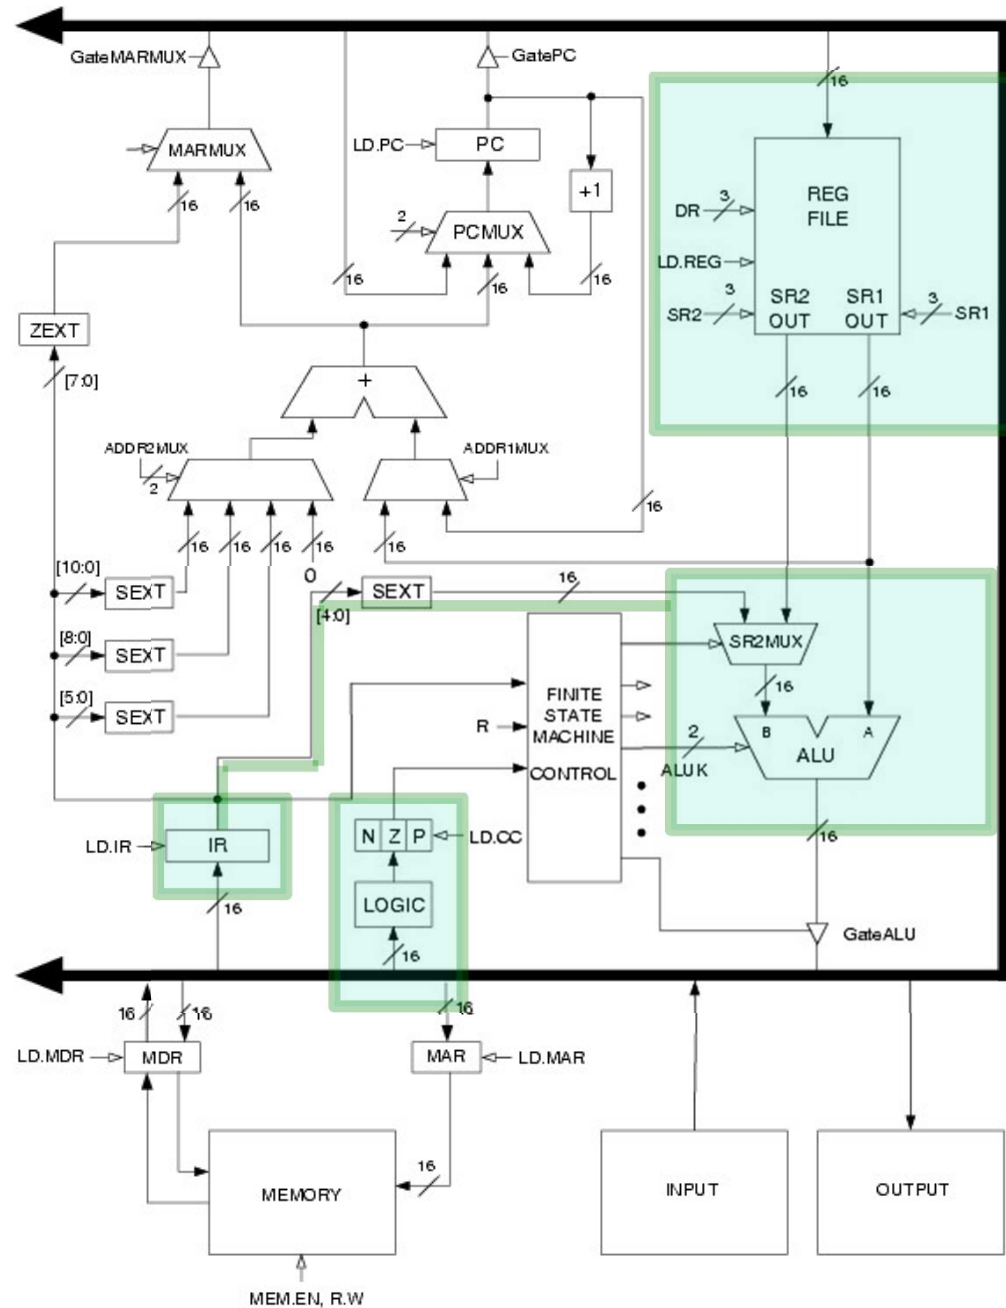
Immediate Values -> x12 for Hex values, decimal values #12



LC-3 Data Path Revisited

Filled arrow
= info to be processed.

Unfilled arrow
= control signal.



Using Operate Instructions

With only ADD, AND, NOT...

- How do we subtract?
- How do we OR?
- How do we copy from one register to another?
- How do we initialize a register to zero?

Control Instructions

Used to alter the sequence of instructions
(by changing the Program Counter)

Conditional Branch

- branch is *taken* if a specified condition is true
 - signed offset is added to PC to yield new PC
- else, the branch is *not taken*
 - PC is not changed, points to the next sequential instruction

Unconditional Branch (or Jump)

- always changes the PC

TRAP

- changes PC to the address of an OS “service routine”
- routine will return control to the next instruction (after TRAP)

Condition Codes

LC-3 has three **condition code** registers:

N -- negative

Z -- zero

P -- positive (greater than zero)

Set by any instruction that writes a value to a register
(ADD, AND, NOT, LD, LDR, LDI, LEA)

Exactly one will be set at all times

- Based on the last instruction that altered a register

Branch Instruction

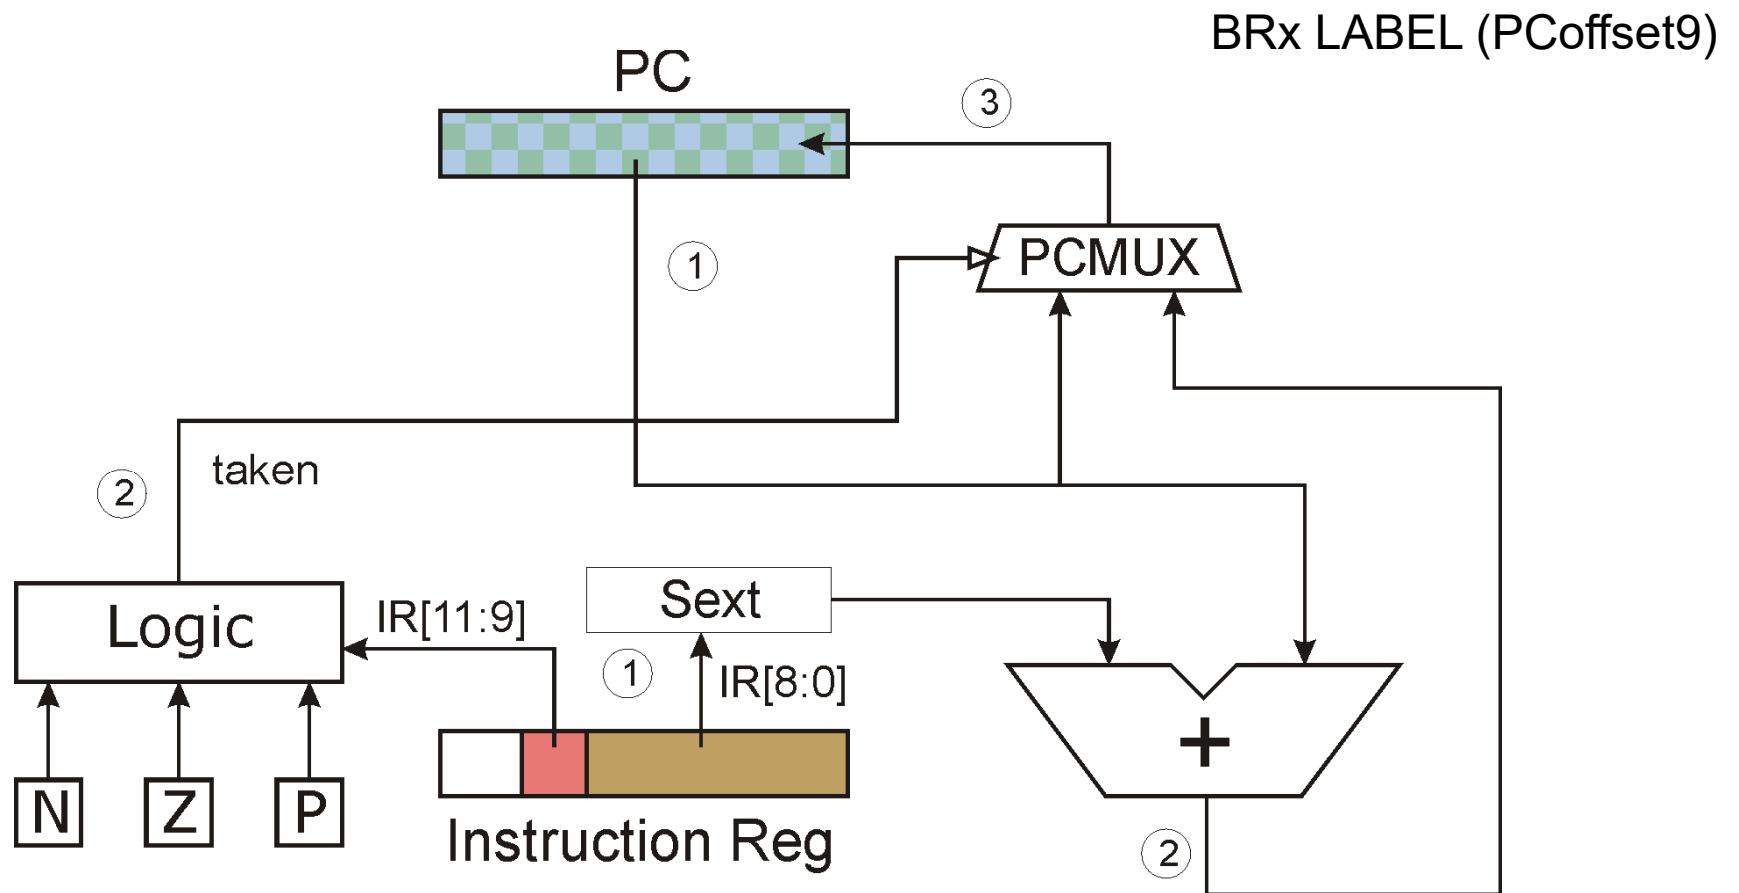

Branch specifies one or more condition codes.

If the set bit is specified, the branch is taken.

- PC-relative addressing:
target address is made by adding signed offset (IR[8:0]) to current PC.
- Note: PC has already been incremented by FETCH stage.
- Note: Target must be within 256 words of BR instruction.

**If the branch is not taken,
the next sequential instruction is executed.**

BR (PC-Relative)

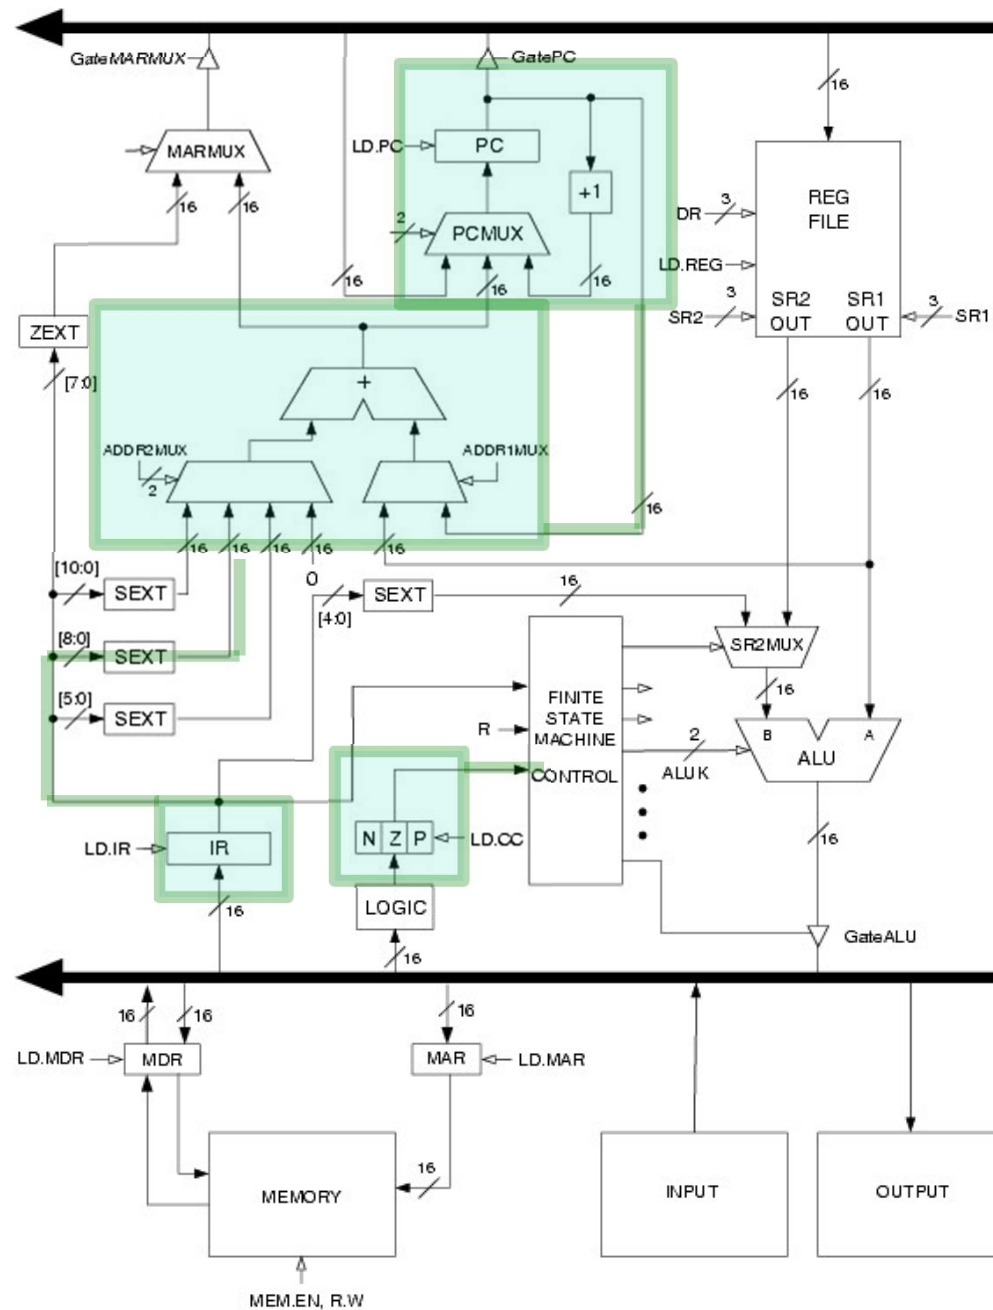


What happens if bits [11:9] are all zero? All one?

LC-3 Data Path Revisited

Filled arrow
= info to be processed.

Unfilled arrow
= control signal.



How are Branches Used?

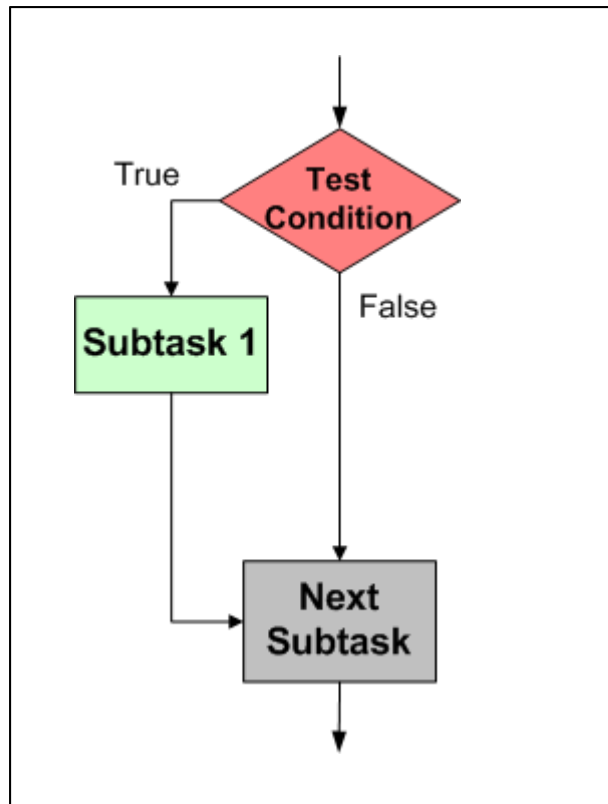
**Alter the flow of the program,
based on some programmer-specified condition,
evaluated at execution time.**

Two basic constructs:

- **Conditional** (If, If-Else)
 - choose whether or not to execute instruction sequence
- **Iterative** (Loops)
 - execute instruction sequence multiple times

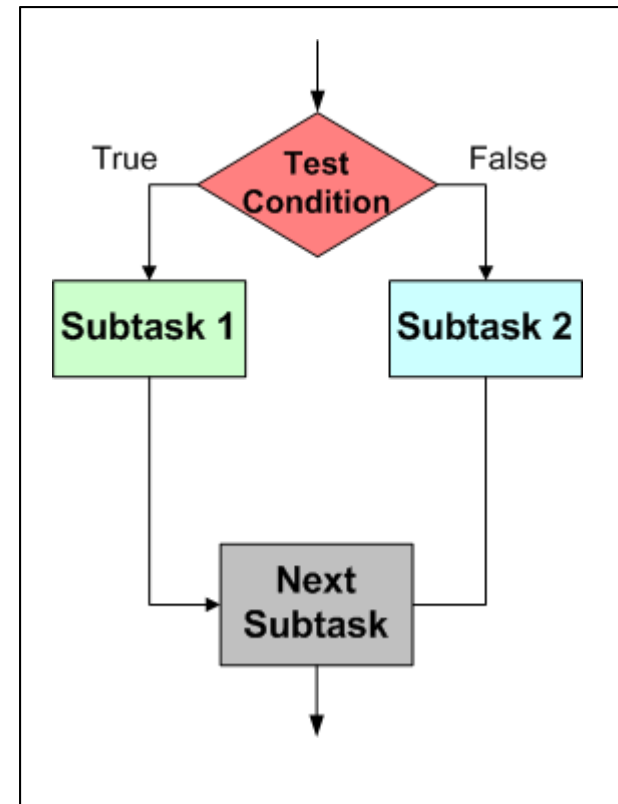
Conditional

If



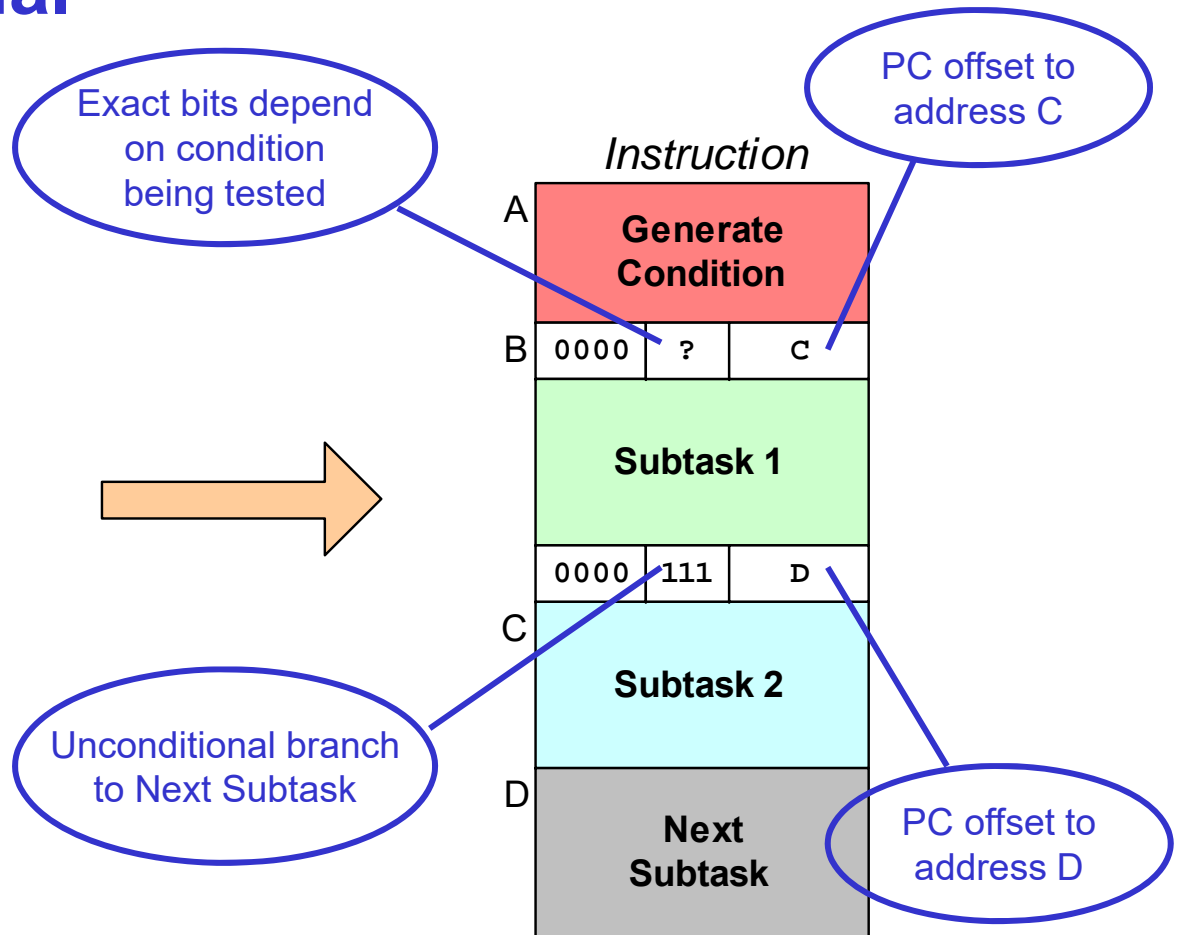
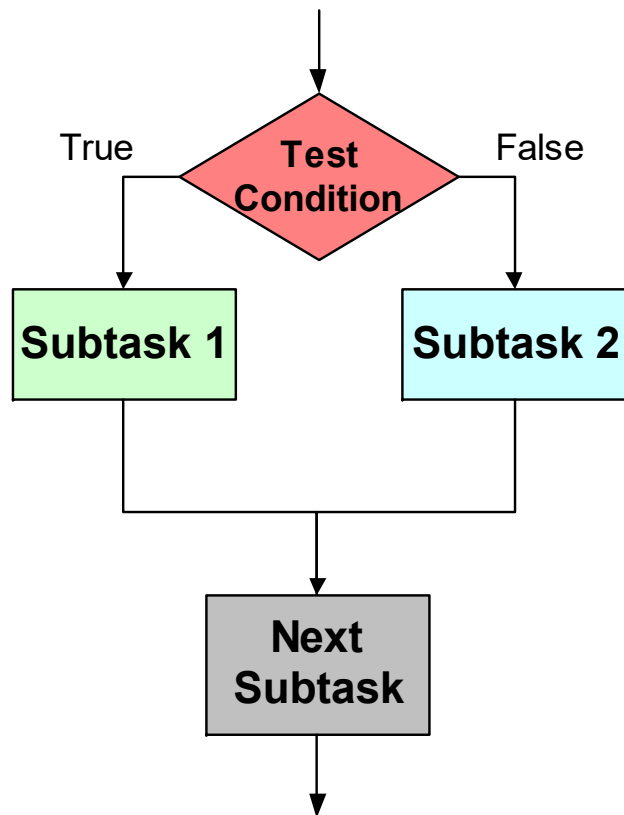
"hammock"

If-Else



"diamond"

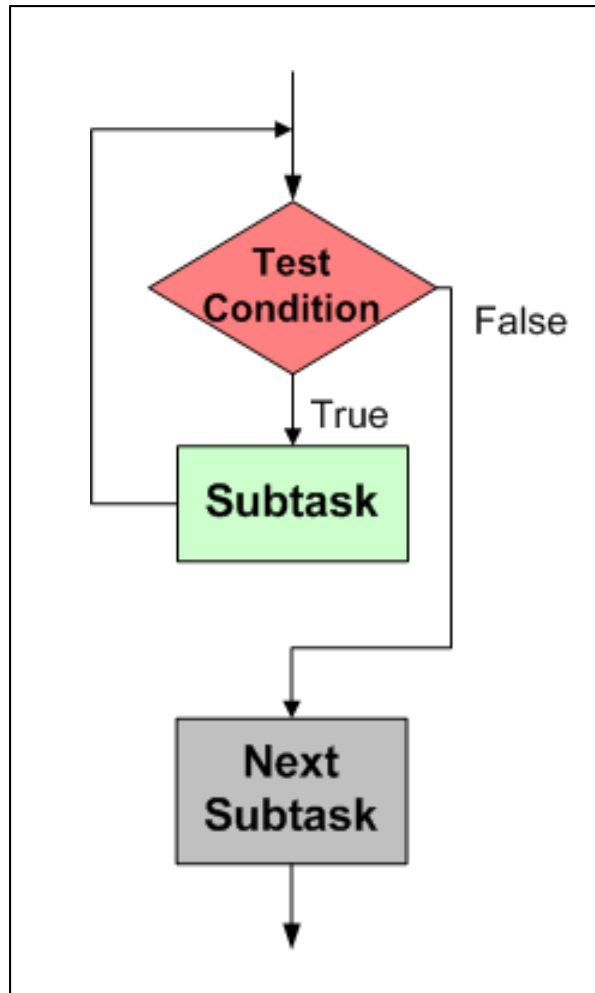
Code for Conditional



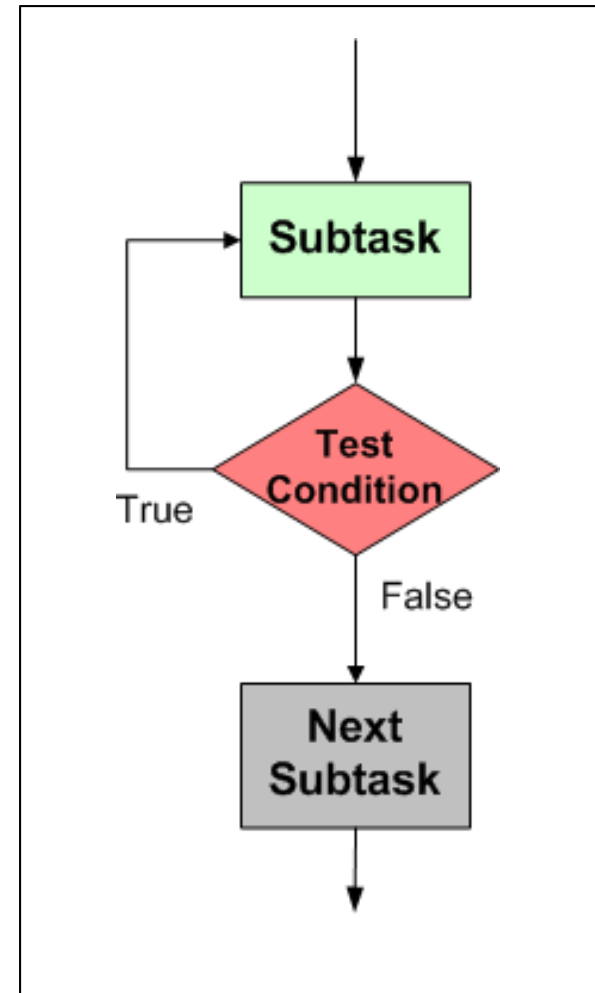
Assumes all addresses are close enough that PC-relative branch can be used.

Iterative

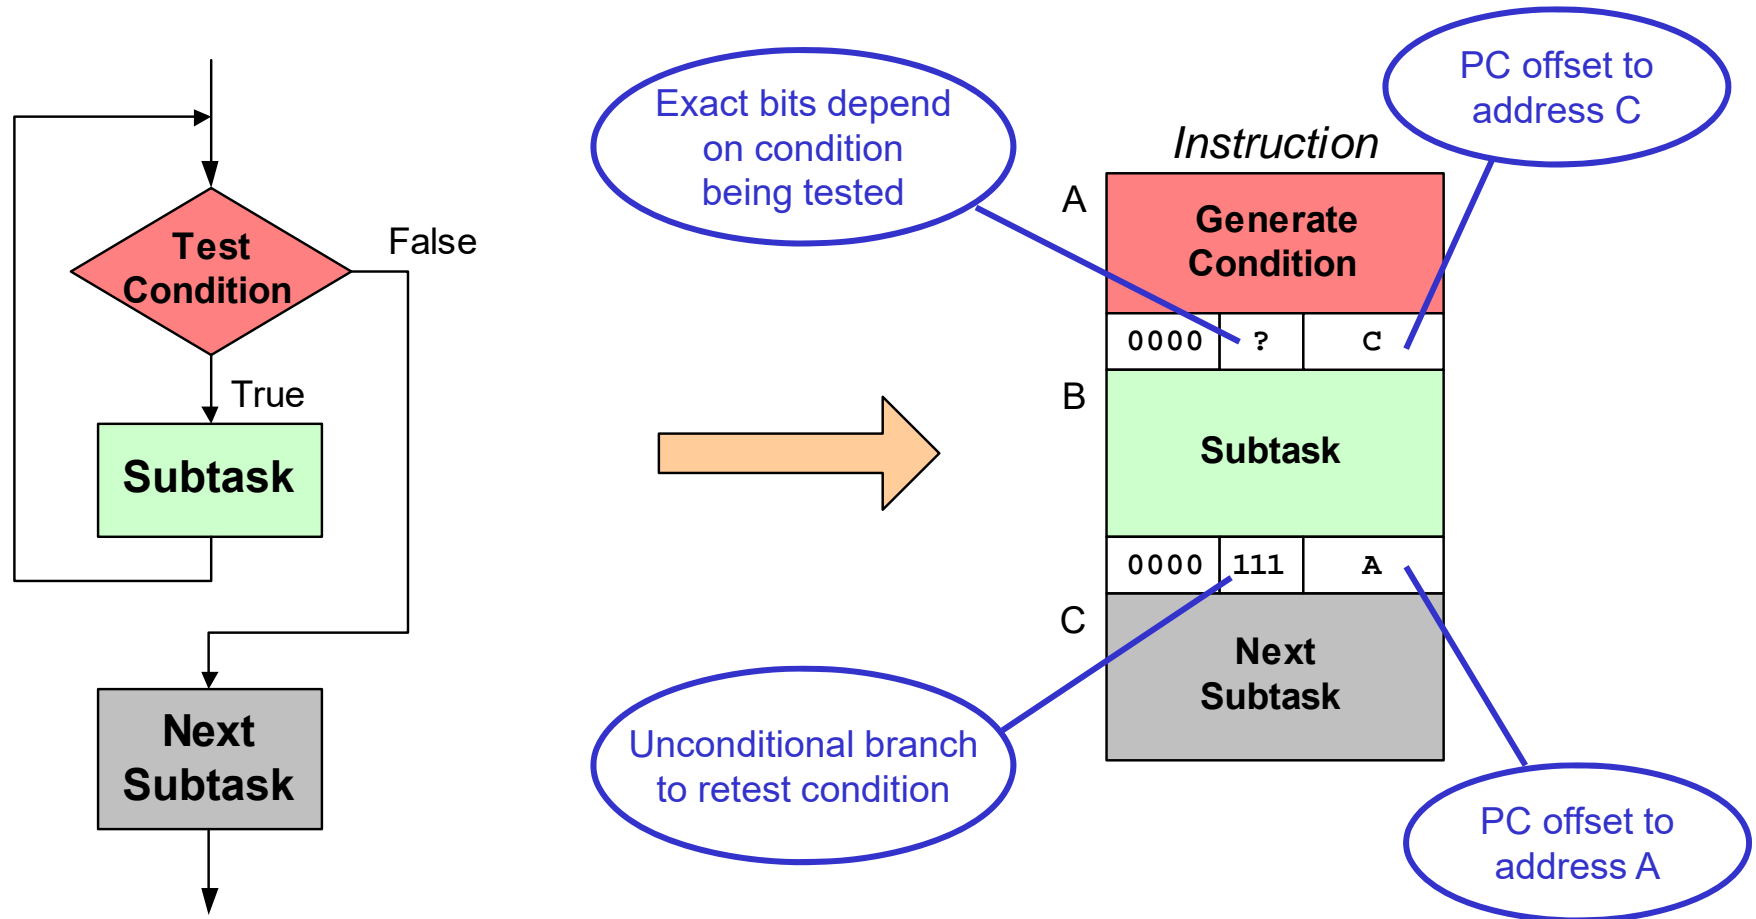
While Loop



Do-While Loop



Code for Iteration



Assumes all addresses are close enough that PC-relative branch can be used.


JMP (Register)

Jump is an unconditional branch -- **always** taken.

- Target address is the contents of a register.
- Allows any target address.


JMP Rbase

	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
JMP	1	1	0	0	0	0	0	Base			0	0	0	0	0	0



TRAP

TRAP trapvect8



Calls a **service routine**, identified by 8-bit “trap vector.”

<i>vector</i>	<i>routine</i>
x23	input a character from the keyboard (IN)
x21	output a character to the monitor (OUT)
x25	halt the program (HALT)

When routine is done,
PC is set to the instruction following TRAP.

(We’ll talk about how this works later.)

Warning: TRAP changes R7.

Data Movement Instructions

Load -- read data from memory to register

- **LD:** PC-relative mode
- **LDR:** base+offset mode
- **LDI:** indirect mode

Store -- write data from register to memory

- **ST:** PC-relative mode
- **STR:** base+offset mode
- **STI:** indirect mode

Load effective address -- compute address, save in register

- **LEA:** immediate mode
- *does not access memory*

PC-Relative Addressing Mode

Want to specify address directly in the instruction

- But an address is 16 bits, and so is an instruction!
- After subtracting 4 bits for opcode and 3 bits for register, we have 9 bits available for address.

Solution:

- Use the 9 bits as a signed offset from the current PC.

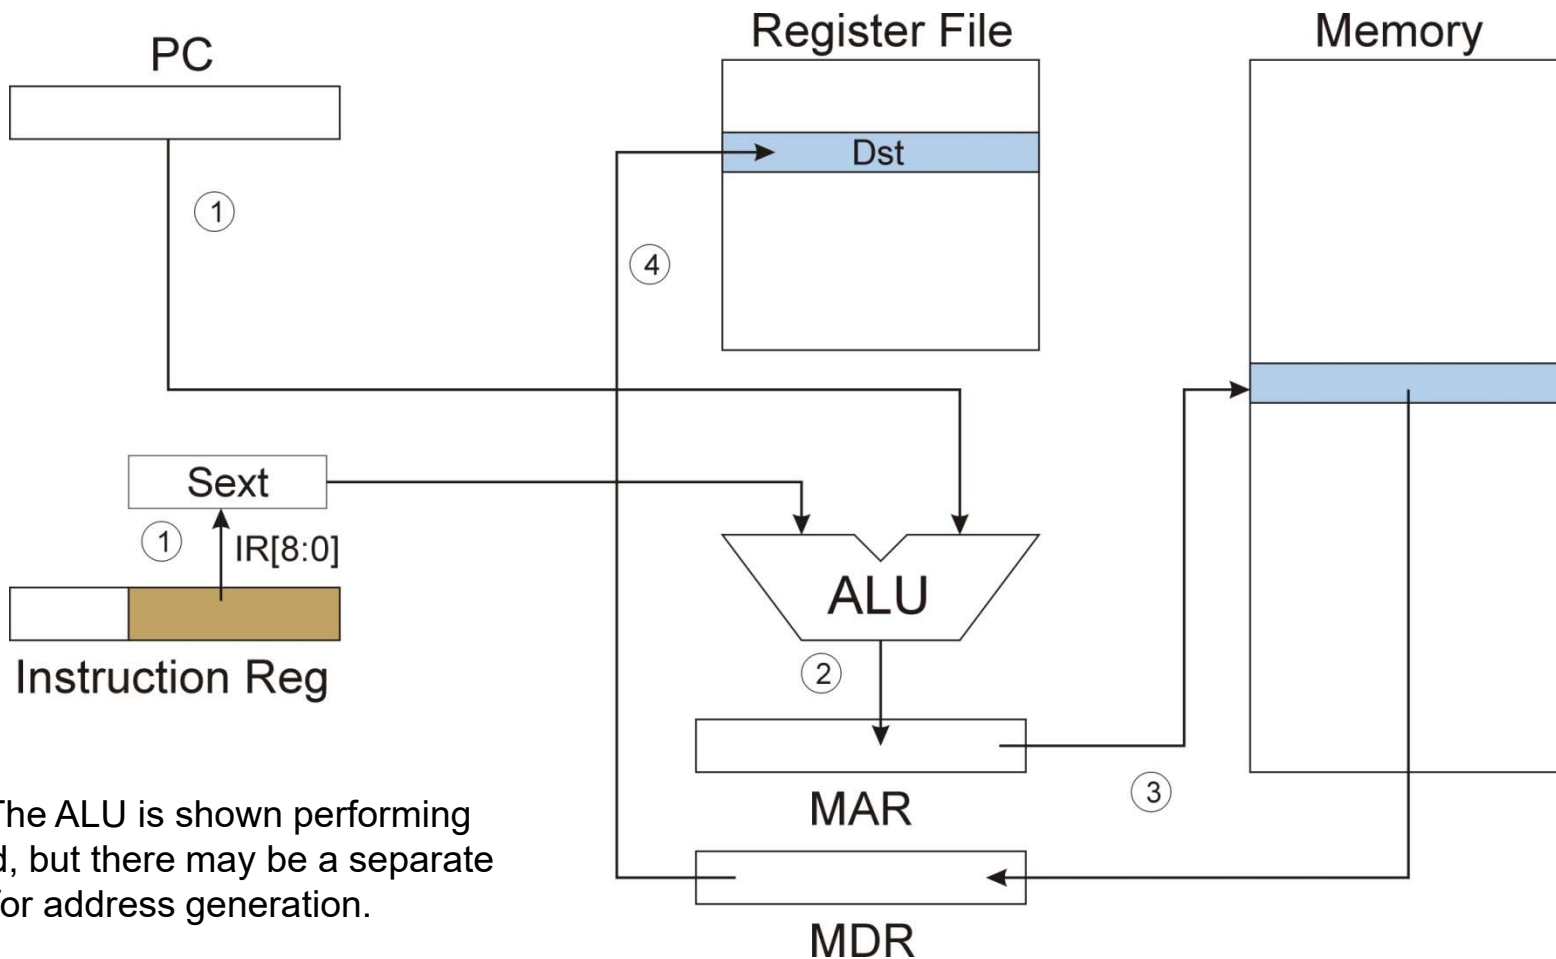

9 bits: $-256 \leq \text{offset} \leq +255$

Can form any address X, such that: $PC - 256 \leq X \leq PC + 255$

**Remember that PC is incremented as part of the FETCH phase;
This is done before the EVALUATE ADDRESS stage.**

LD (PC-Relative)

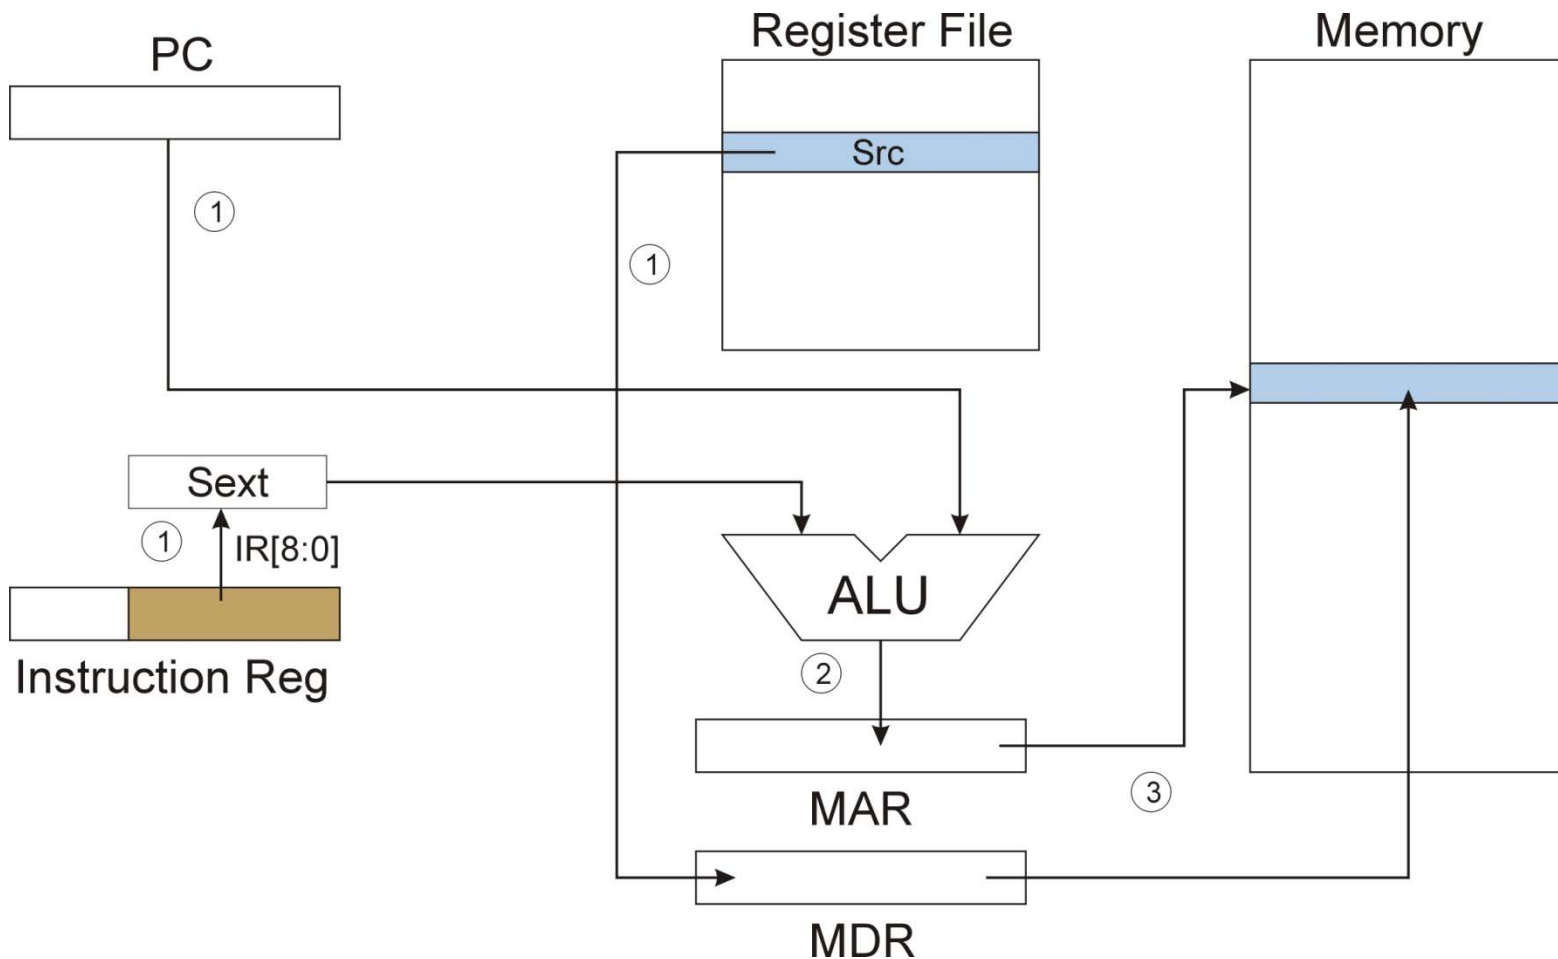

LD Dst, LABEL (PCoffset9)



Note: The ALU is shown performing the add, but there may be a separate adder for address generation.

ST (PC-Relative)

ST Src, LABEL (PCOffset9)



Base + Offset Addressing Mode

With PC-relative mode, can only address data within 256 words of the instruction.

- What about the rest of memory?

Solution #1:

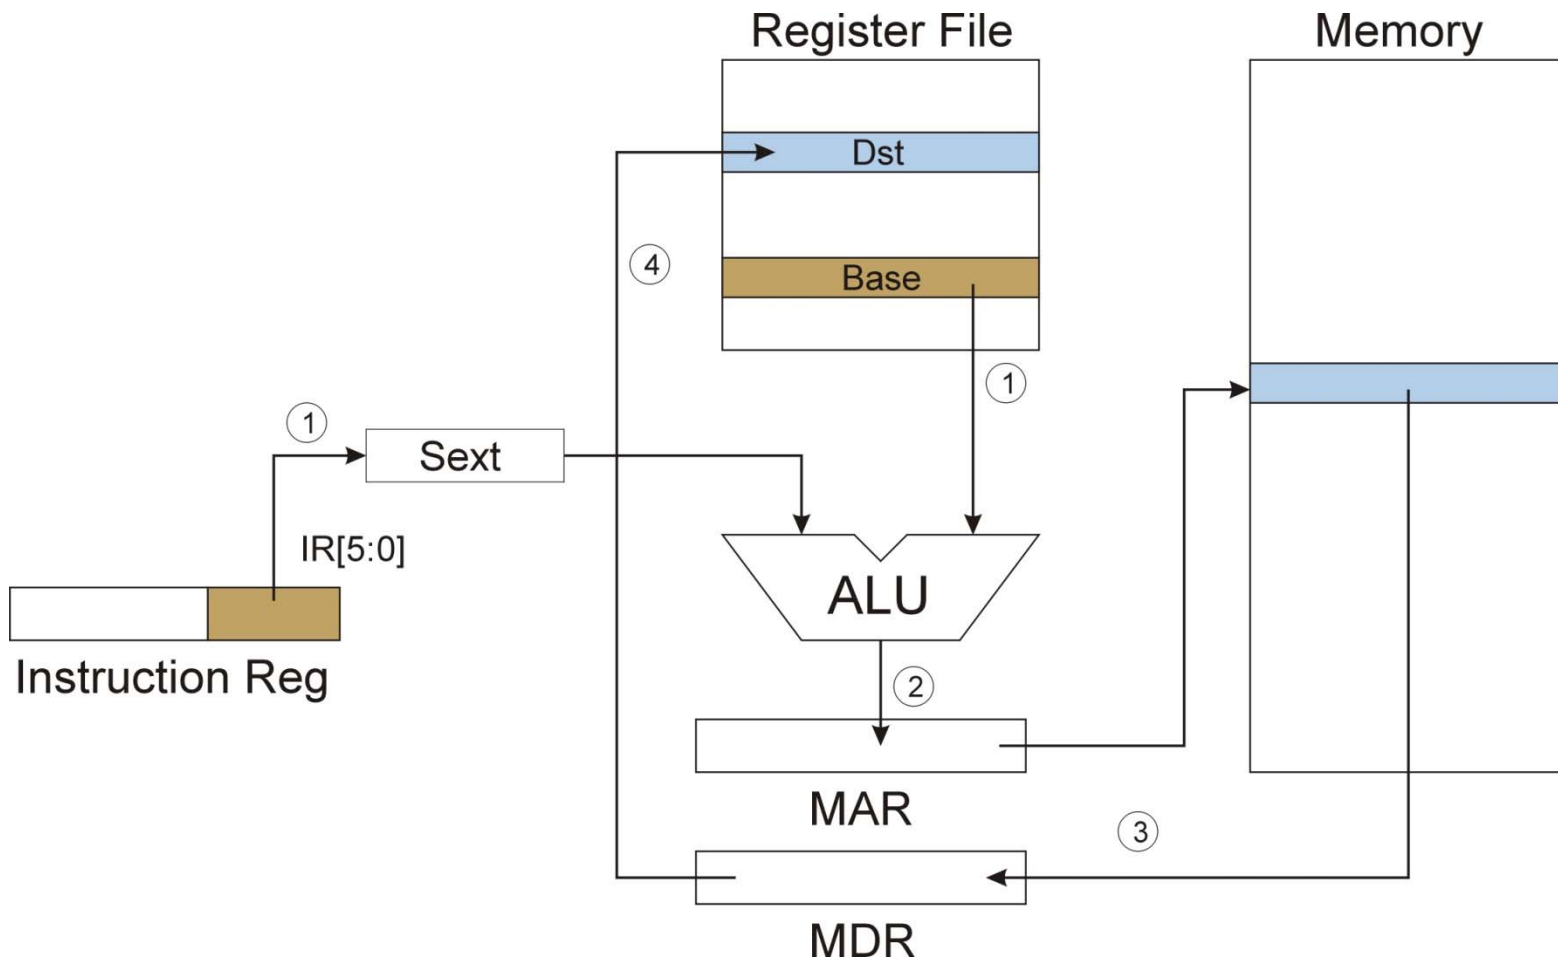

- Use the value in a register to generate a full 16-bit address.

4 bits for opcode, 3 for src/dest register,
3 bits for **base** register -- remaining 6 bits are used
as a **signed offset**.

- Offset is *sign-extended* before adding to base register.

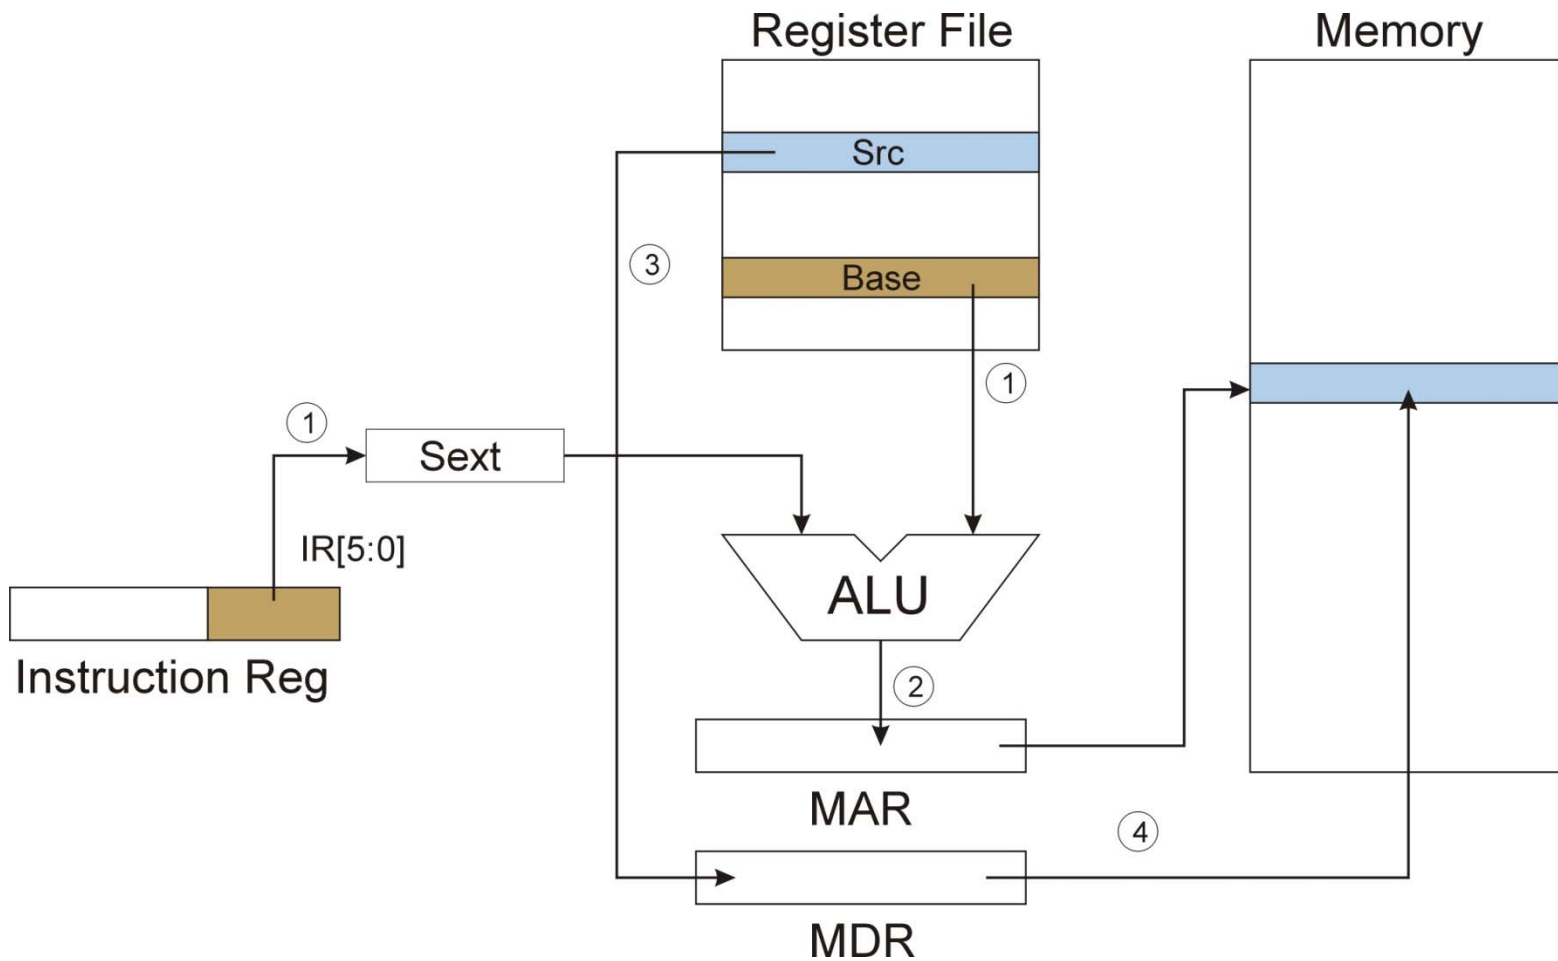

LDR (Base+Offset)

LDR Dst, BaseR LABEL (offset6)



STR (Base+Offset)

STR Src, BaseR LABEL (offset6)



Indirect Addressing Mode

With PC-relative mode, can only address data within 256 words of the instruction.

- What about the rest of memory?

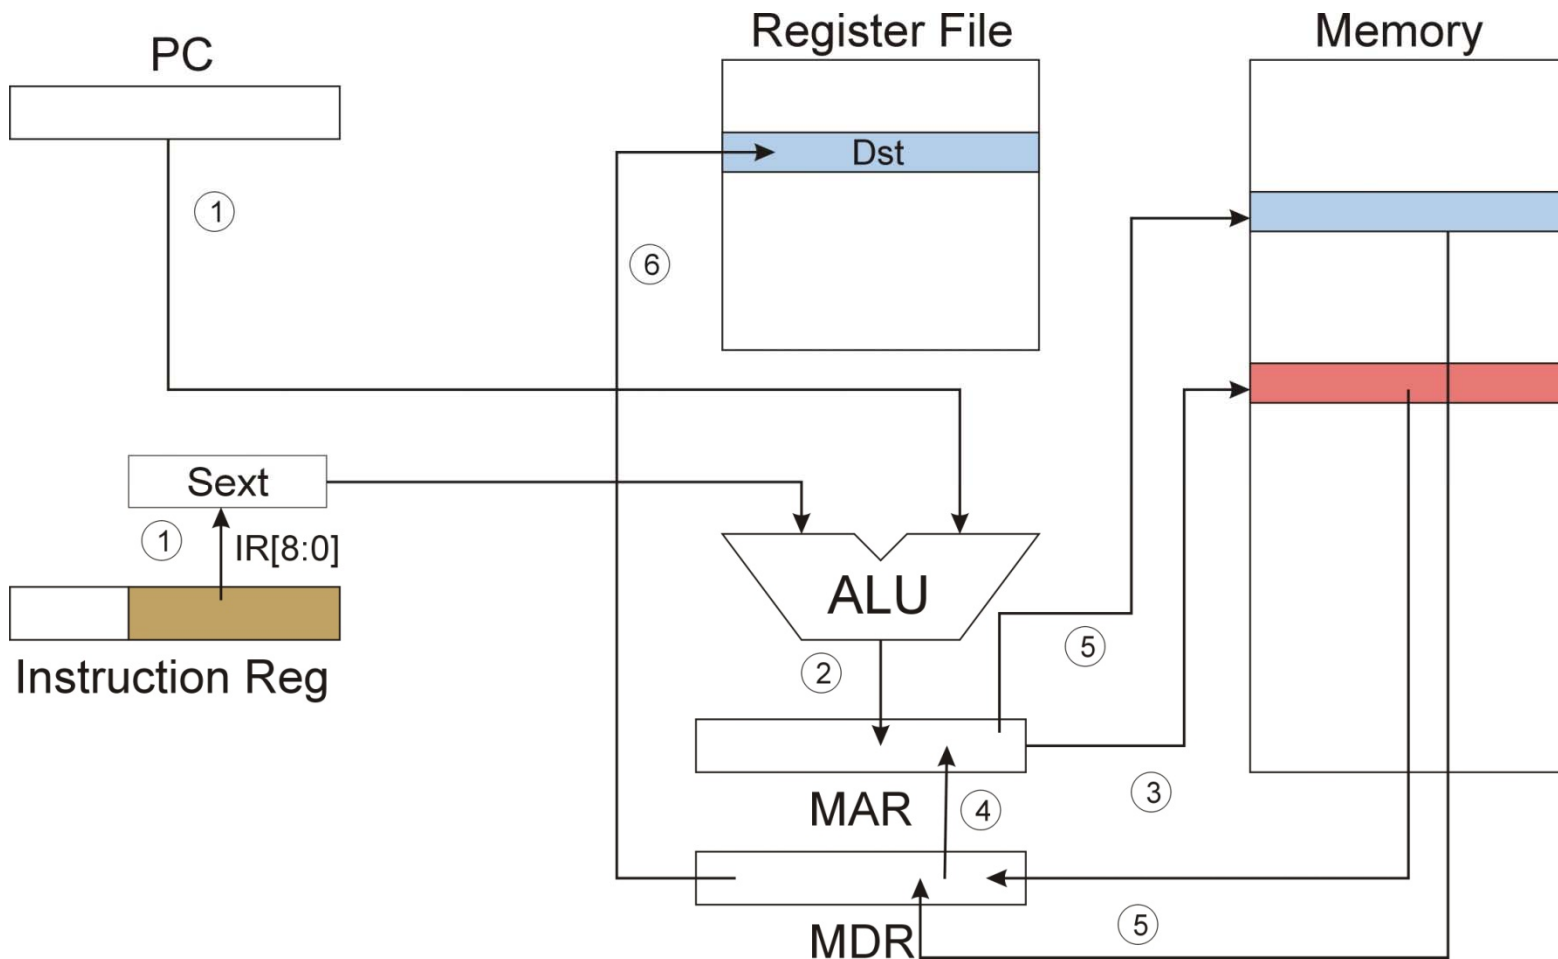

Solution #3:

- **Read address from memory location, then load/store to that address.**

First address is generated from PC and IR (just like PC-relative addressing), then content of that address is used as target for load/store.

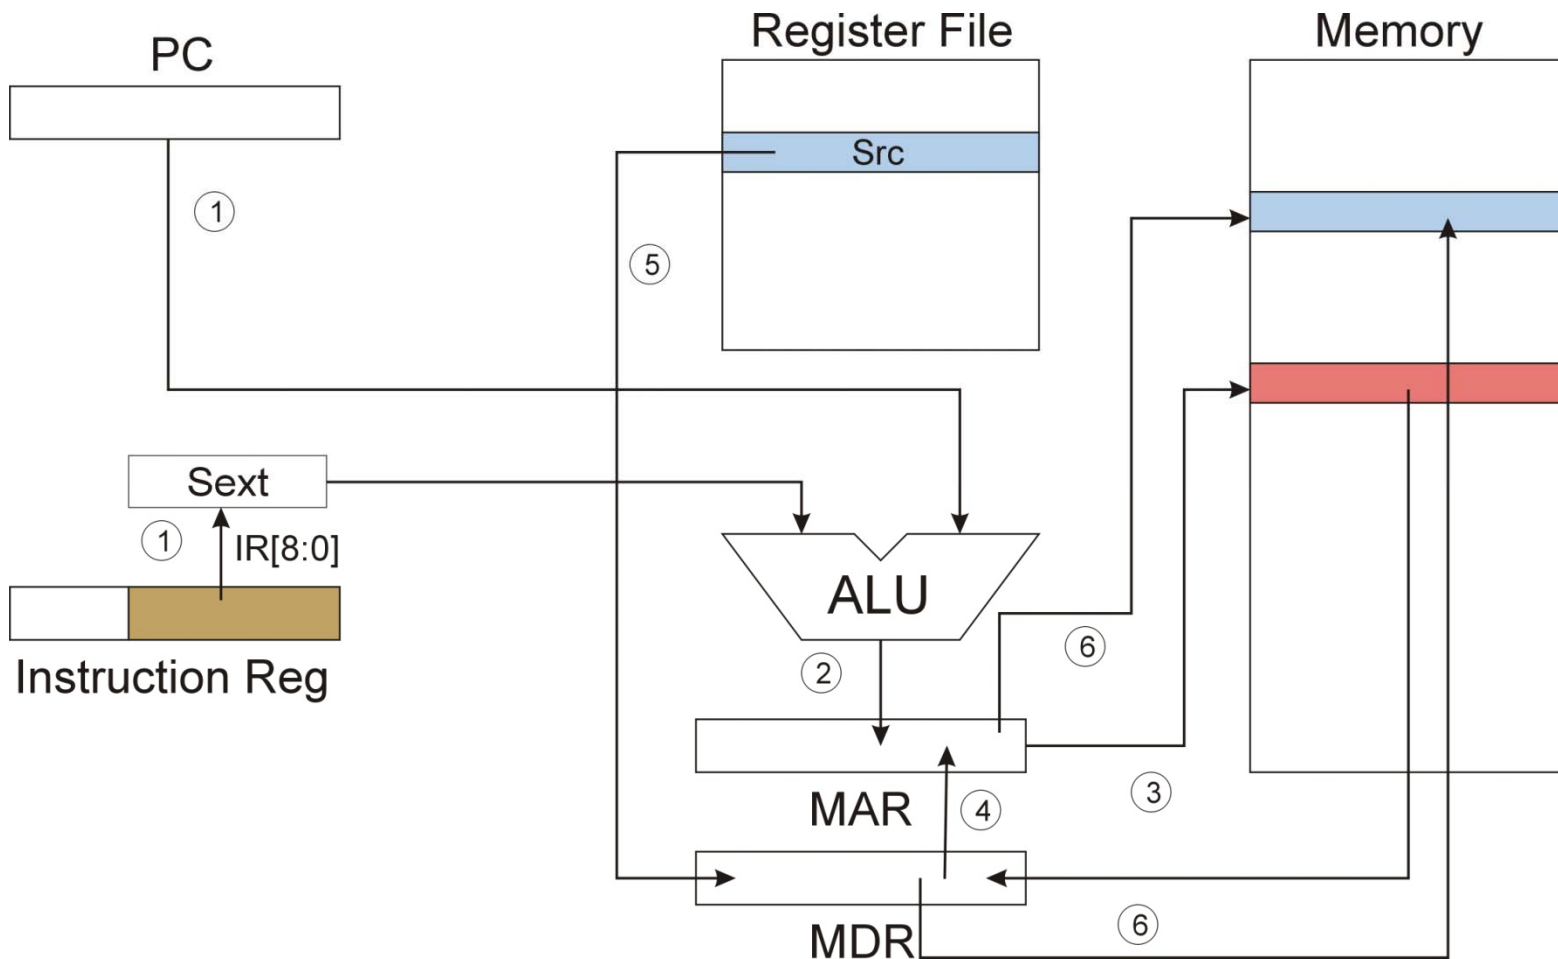

LDI (Indirect)

LDI Dst, LABEL (PCoffset9)



STI (Indirect)

STI Src, LABEL (PCoffset9)



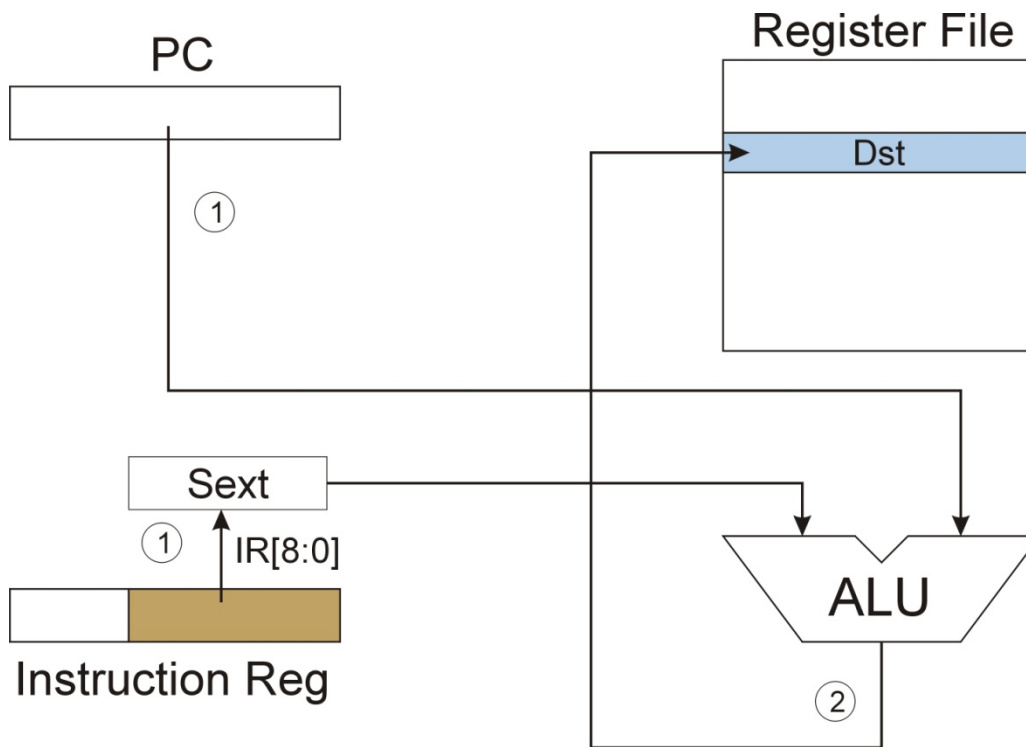

Load Effective Address

Computes address like PC-relative (PC plus signed offset) and **stores the result into a register.**

Note: The address is stored in the register, not the contents of the memory location.

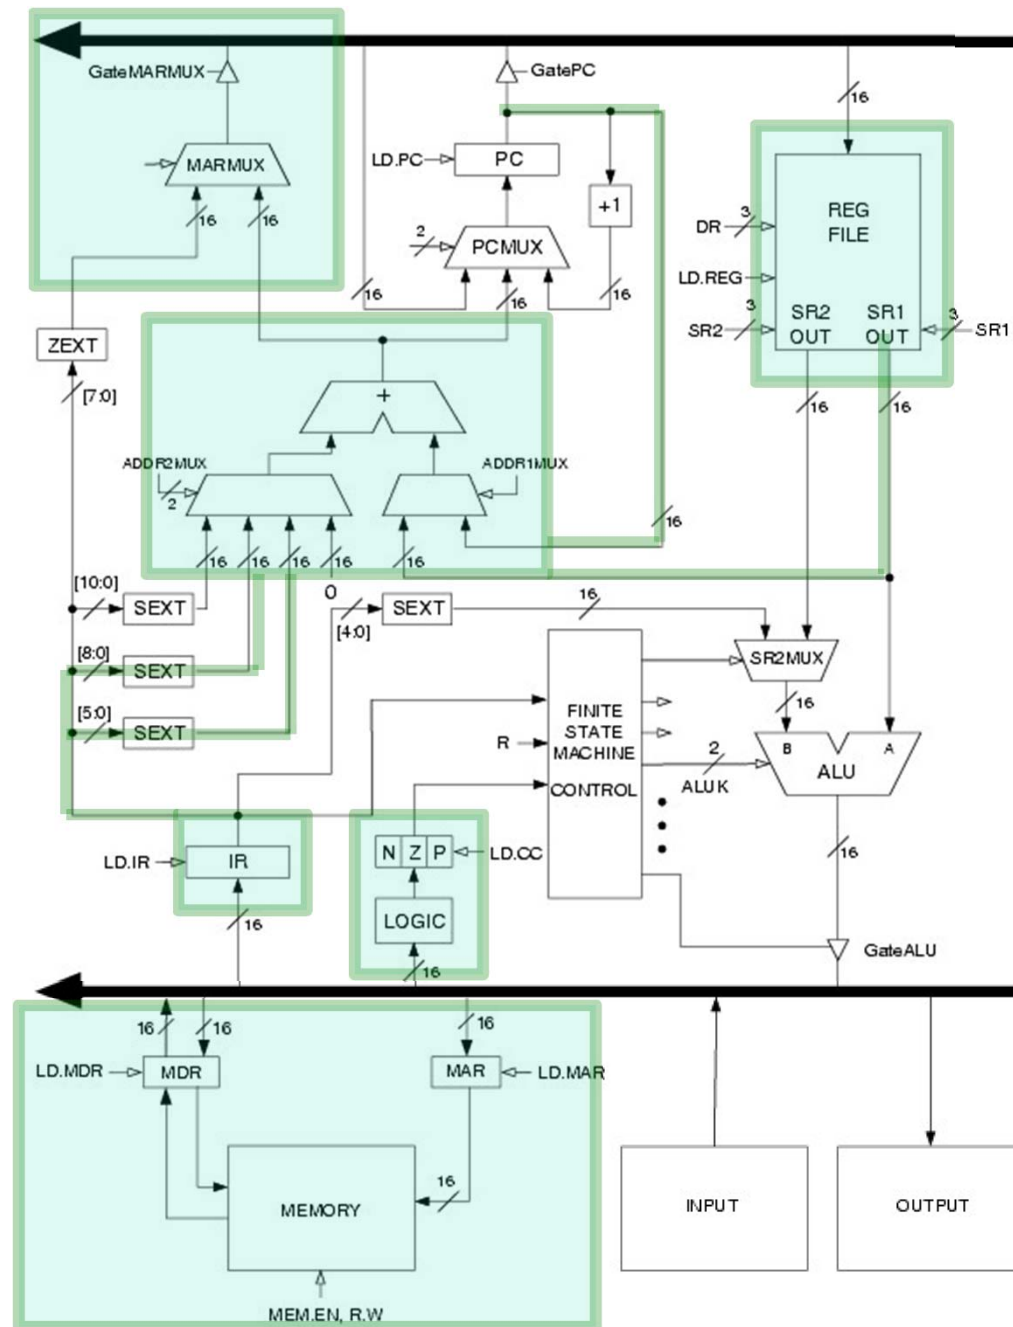
LEA (Immediate)

LEA Dst LABEL (PCoffset9)



LC-3 Data Path Revisited

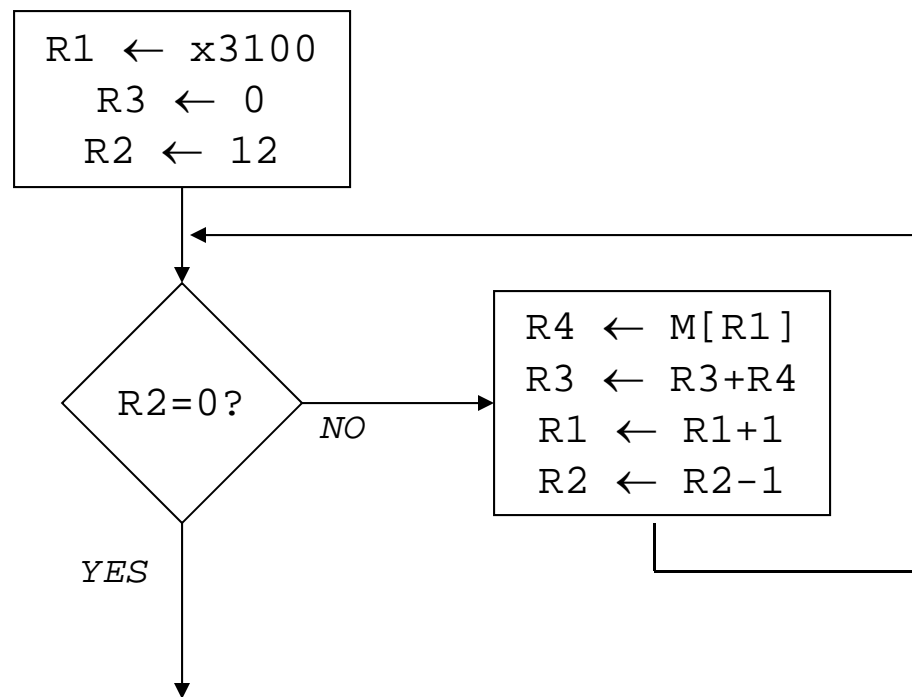
Filled arrow
= info to be processed.
Unfilled arrow
= control signal.



Using Branch and Load Instructions

Compute sum of 12 integers.

Numbers start at location x3100. Program starts at location x3000.



Sample Program

Address	Instruction														Comments
x3000	1	1	1	0	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<i>R1 ← x3100 (PC+0xFF)</i>
x3001	0	1	0	1	<u>0</u>	<u>1</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>0</u>	<i>R3 ← 0</i>
x3002	0	1	0	1	<u>0</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>0</u>	<i>R2 ← 0</i>
x3003	0	0	0	1	<u>0</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>1</u>	<i>R2 ← 12</i>
x3004	0	0	0	0	<u>0</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>1</u>	<i>If Z, goto x300A (PC+5)</i>
x3005	0	1	1	0	<u>1</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<i>Load next value to R4</i>
x3006	0	0	0	1	<u>0</u>	<u>1</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<i>Add to R3</i>
x3007	0	0	0	1	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>1</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>0</u>	<i>Increment R1 (pointer)</i>
x3008	0	0	0	1	<u>0</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<i>Decrement R2 (counter)</i>
x3009	0	0	0	0	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>0</u>	<i>Goto x3004 (PC-6)</i>

Another Example

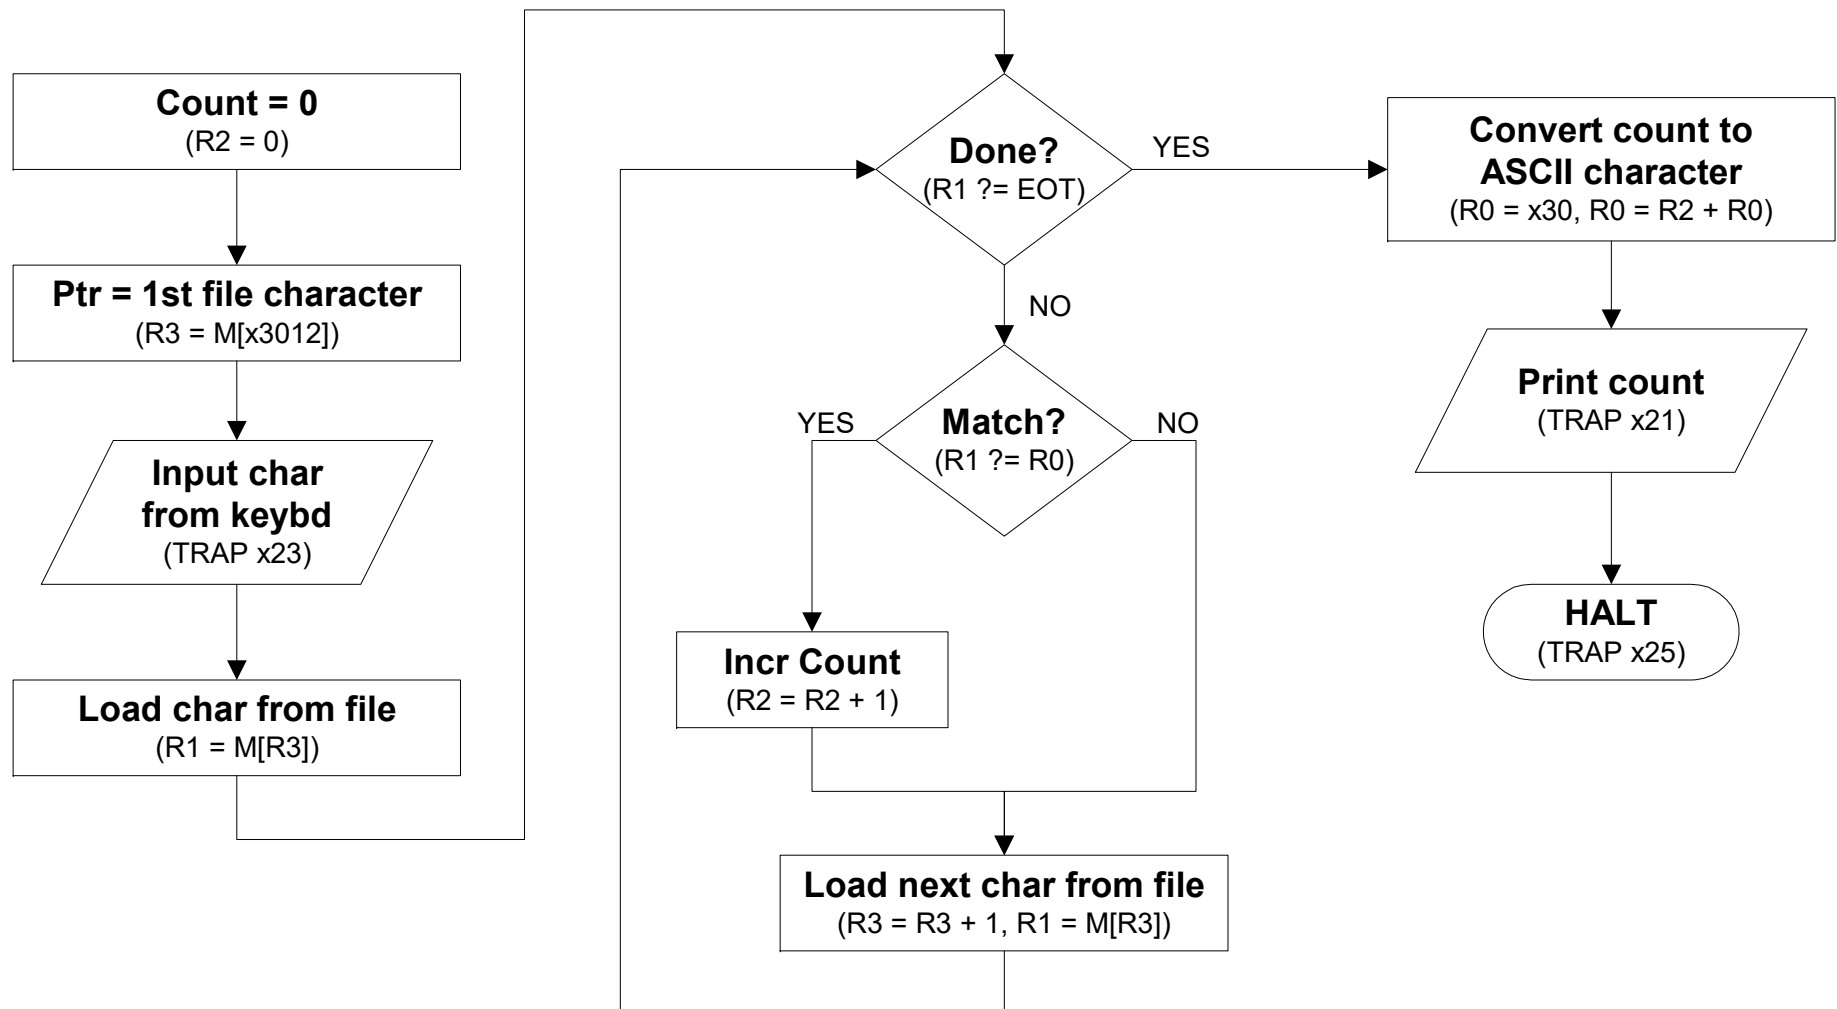
Count the occurrences of a character in a memory region

- Program begins at location x3000
- Read character from keyboard
- Load each character from a “file”
 - File is a sequence of memory locations
 - Starting address of file is stored in the memory location immediately after the program
- If file character equals input character, increment counter
- End of file is indicated by a special ASCII value: **EOT (x04)**
- At the end, print the number of characters and halt
(assume there will be less than 10 occurrences of the character)

A special character used to indicate the end of a sequence is often called a **sentinel.**

- Useful when you don't know ahead of time how many times to execute a loop.

Flow Chart



Program (1 of 2)

Address	Instruction														Comments	
x3000	0	1	0	1	0	1	0	0	1	0	1	0	0	0	0	$R2 \leftarrow 0$ (counter)
x3001	0	0	1	0	0	1	1	0	0	0	0	1	0	0	0	$R3 \leftarrow M[x3102]$ (ptr)
x3002	1	1	1	1	0	0	0	0	0	0	1	0	0	0	1	Input to R0 (TRAP x23)
x3003	0	1	1	0	0	0	1	0	1	1	0	0	0	0	0	$R1 \leftarrow M[R3]$
x3004	0	0	0	1	1	0	0	0	0	1	1	1	1	1	0	$R4 \leftarrow R1 - 4$ (EOT)
x3005	0	0	0	0	0	1	0	0	0	0	0	0	1	0	0	If Z, goto x300E
x3006	1	0	0	1	0	0	1	0	0	1	1	1	1	1	1	$R1 \leftarrow NOT\ R1$
x3007	0	0	0	1	0	0	1	0	0	1	1	0	0	0	0	$R1 \leftarrow R1 + 1$
x3008	0	0	0	1	0	0	1	0	0	1	0	0	0	0	0	$R1 \leftarrow R1 + R0$
x3009	0	0	0	0	1	0	1	0	0	0	0	0	0	0	0	If N or P, goto x300B

Program (2 of 2)

Address	Instruction														Comments		
x300A	0	0	0	1	<u>0</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>1</u>	<i>R2 ← R2 + 1</i>
x300B	0	0	0	1	<u>0</u>	<u>1</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>1</u>	1	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>1</u>	<i>R3 ← R3 + 1</i>
x300C	0	1	1	0	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<i>R1 ← M[R3]</i>
x300D	0	0	0	0	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>1</u>	<u>0</u>	<i>Goto x3004</i>
x300E	0	0	1	0	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>0</u>	<i>R0 ← M[x3013]</i>
x300F	0	0	0	1	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<i>R0 ← R0 + R2</i>
x3010	1	1	1	1	0	0	0	0	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>0</u>	<u>1</u>	<i>Print R0 (TRAP x21)</i>
x3011	1	1	1	1	0	0	0	0	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>1</u>	<i>HALT (TRAP x25)</i>
X3012	Starting Address of File																
x3013	0	0	0	0	0	0	0	0	0	0	1	1	0	0	0	0	<i>ASCII x30 ('0')</i>