

Chapter 8

I/O

I/O: Connecting to Outside World

So far, we've learned how to:

- **compute with values in registers**
- **load data from memory to registers**
- **store data from registers to memory**

But where does data in memory come from?

And how does data get out of the system so that humans can use it?

I/O: Connecting to the Outside World

Types of I/O devices characterized by:

- **behavior:** input, output, storage
 - input: keyboard, motion detector, network interface
 - output: monitor, printer, network interface
 - storage: disk, CD-ROM
- **data rate:** how fast can data be transferred?
 - keyboard: 100 bytes/sec
 - disk: 30 MB/s
 - network: 1 Mb/s - 1 Gb/s

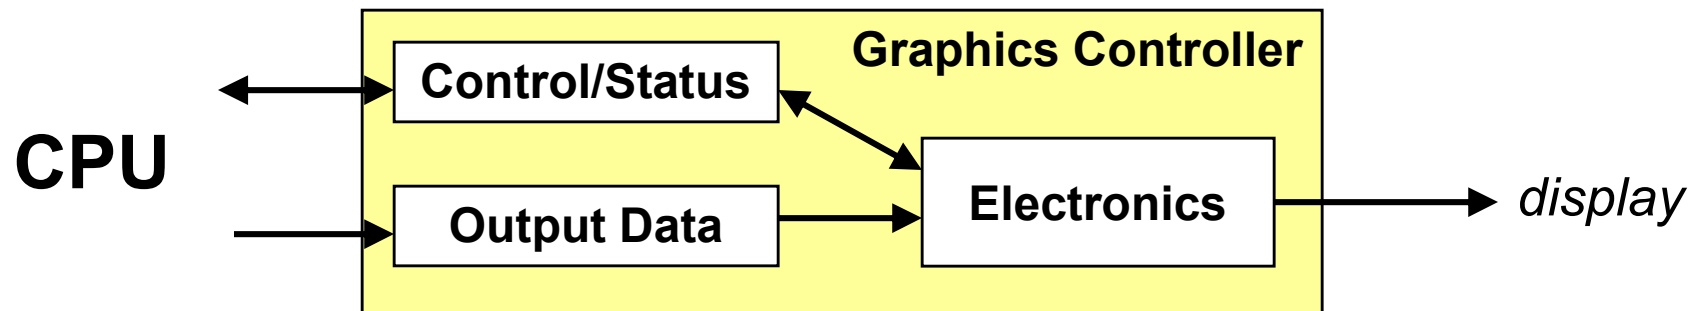
I/O Controller

Control/Status Registers

- CPU tells device what to do -- write to control register
- CPU checks whether task is done -- read status register

Data Registers

- CPU transfers data to/from device



Device electronics

- performs actual operation
 - pixels to screen, bits to/from disk, characters from keyboard

Programming Interface

How are device registers identified?

- **Memory-mapped** vs. **special instructions**

How is timing of transfer managed?

- **Asynchronous** vs. **synchronous**


Who controls transfer?

- CPU (**polling**) vs. device (**interrupts**)

Memory-Mapped vs. I/O Instructions


Instructions

- designate opcode(s) for I/O
- register and operation encoded in instruction



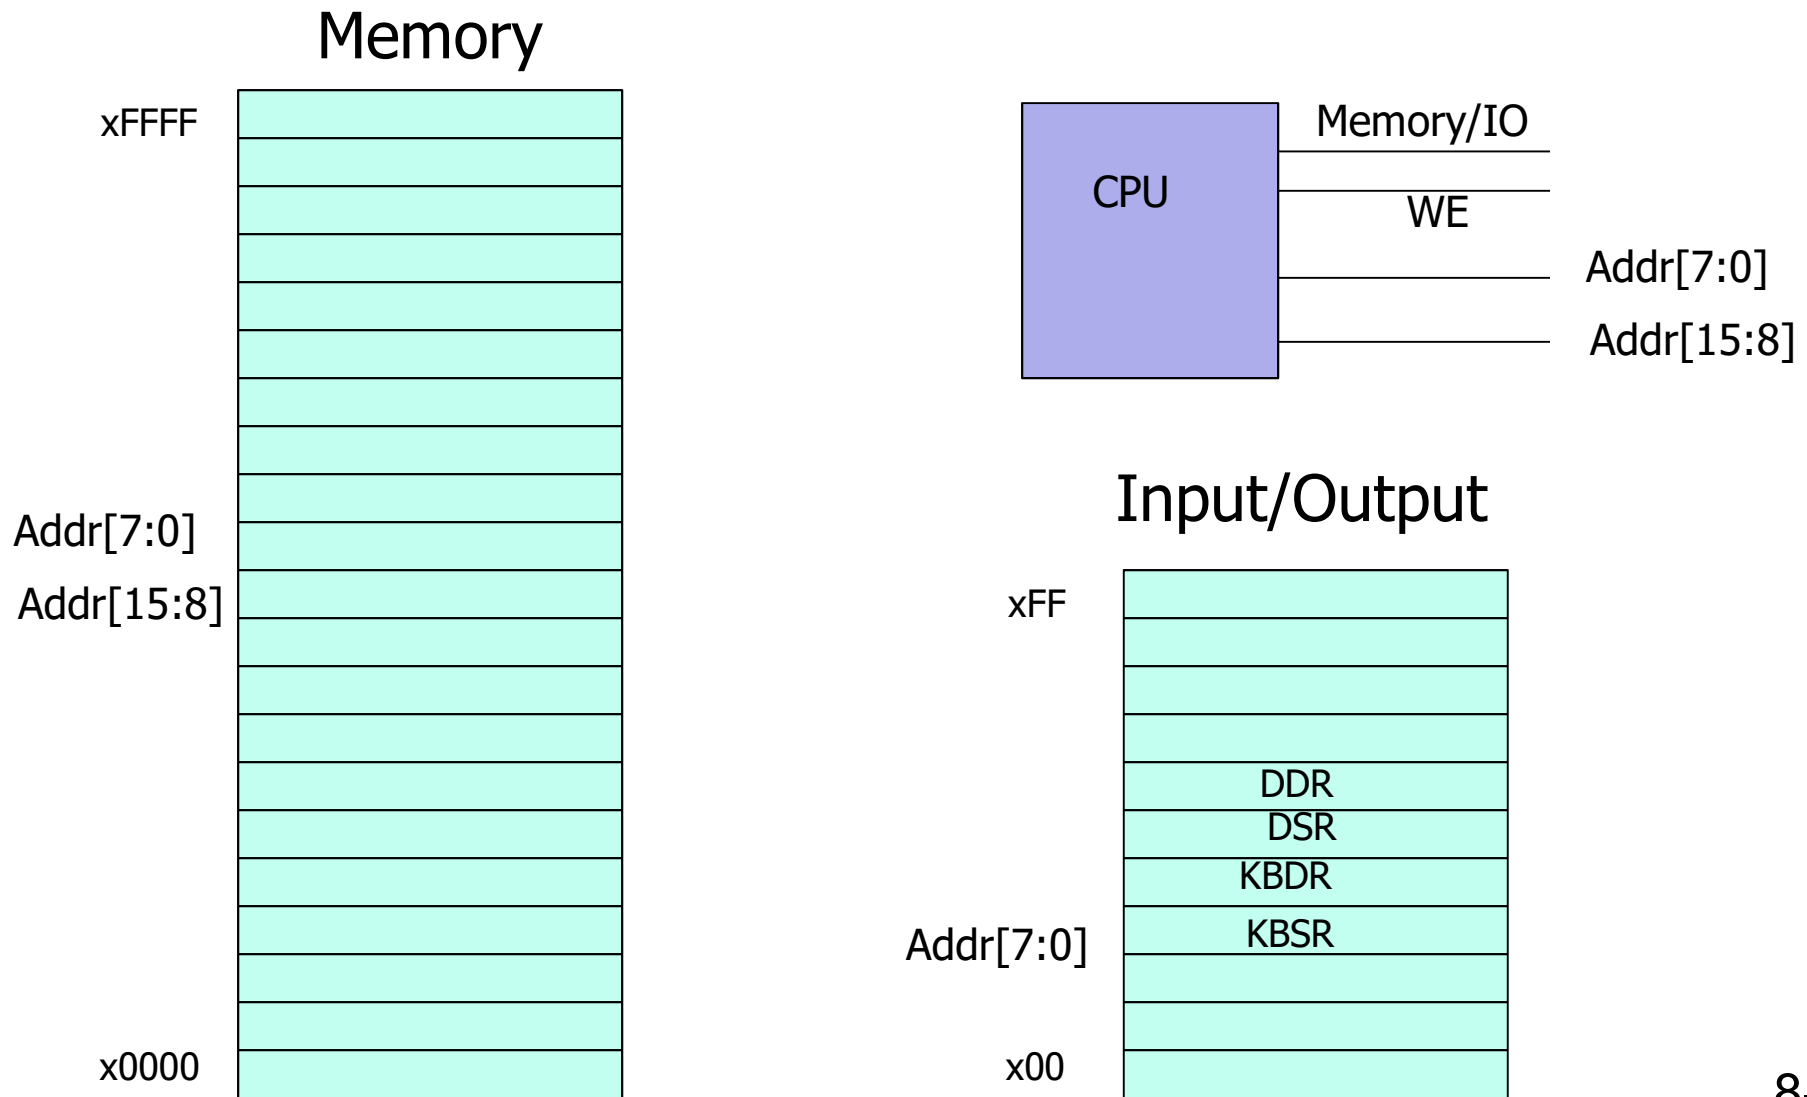
Memory-mapped

- assign a memory address to each device register
- use data movement instructions (LD/ST) for control and data transfer



Memory-Mapped vs. I/O Instructions

I/O-mapped



Transfer Timing

I/O events generally happen much slower than CPU cycles.

Synchronous

- data supplied at a fixed, predictable rate
- CPU reads/writes every X cycles

Asynchronous

- data rate less predictable
- CPU must synchronize with device, so that it doesn't miss data or write too quickly

Transfer Control

Who determines when the next data transfer occurs?

Polling

- CPU keeps checking status register until new data arrives OR device ready for next data
- “Are we there yet? Are we there yet? Are we there yet?”

Interrupts

- Device sends a special signal to CPU when new data arrives OR device ready for next data
- CPU can be performing other tasks instead of polling device.
- “Wake me when we get there.”

LC-3

Memory-mapped I/O (Table A.3)

<i>Location</i>	<i>I/O Register</i>	<i>Function</i>
xFE00	Keyboard Status Reg (KBSR)	Bit [15] is one when keyboard has received a new character.
xFE02	Keyboard Data Reg (KBDR)	Bits [7:0] contain the last character typed on keyboard.
xFE04	Display Status Register (DSR)	Bit [15] is one when device ready to display another char on screen.
xFE06	Display Data Register (DDR)	Character written to bits [7:0] will be displayed on screen.

Asynchronous devices

- synchronized through status registers

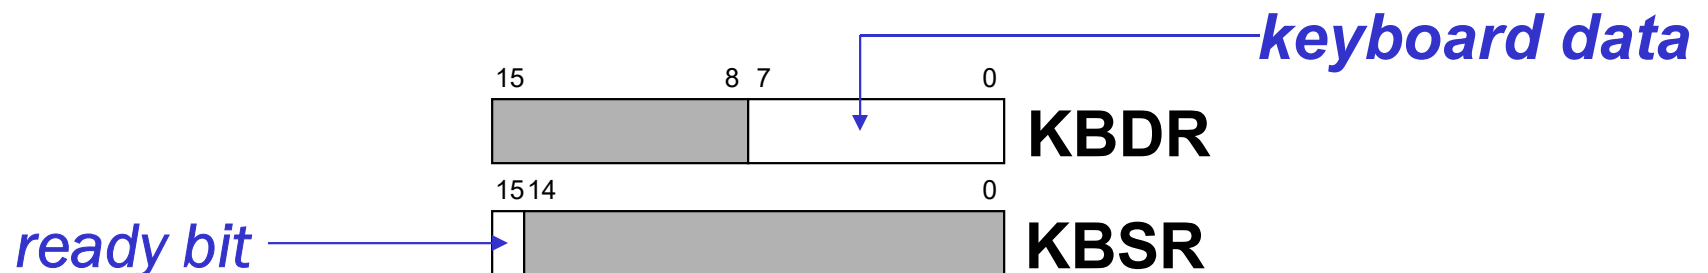
Polling and Interrupts

- the details of interrupts will be discussed in Chapter 10

Input from Keyboard

When a character is typed:

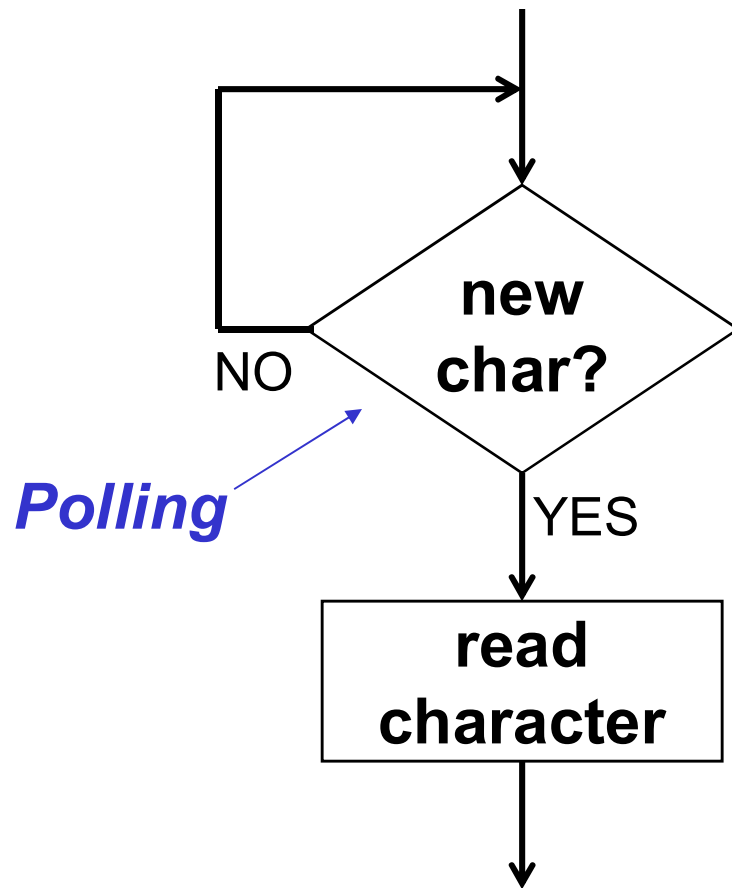
- its ASCII code is placed in bits [7:0] of KBDR (bits [15:8] are always zero)
- the “ready bit” (KBSR[15]) is set to one
- keyboard is disabled -- any typed characters will be ignored



When KBDR is read:

- KBSR[15] is set to zero
- keyboard is enabled

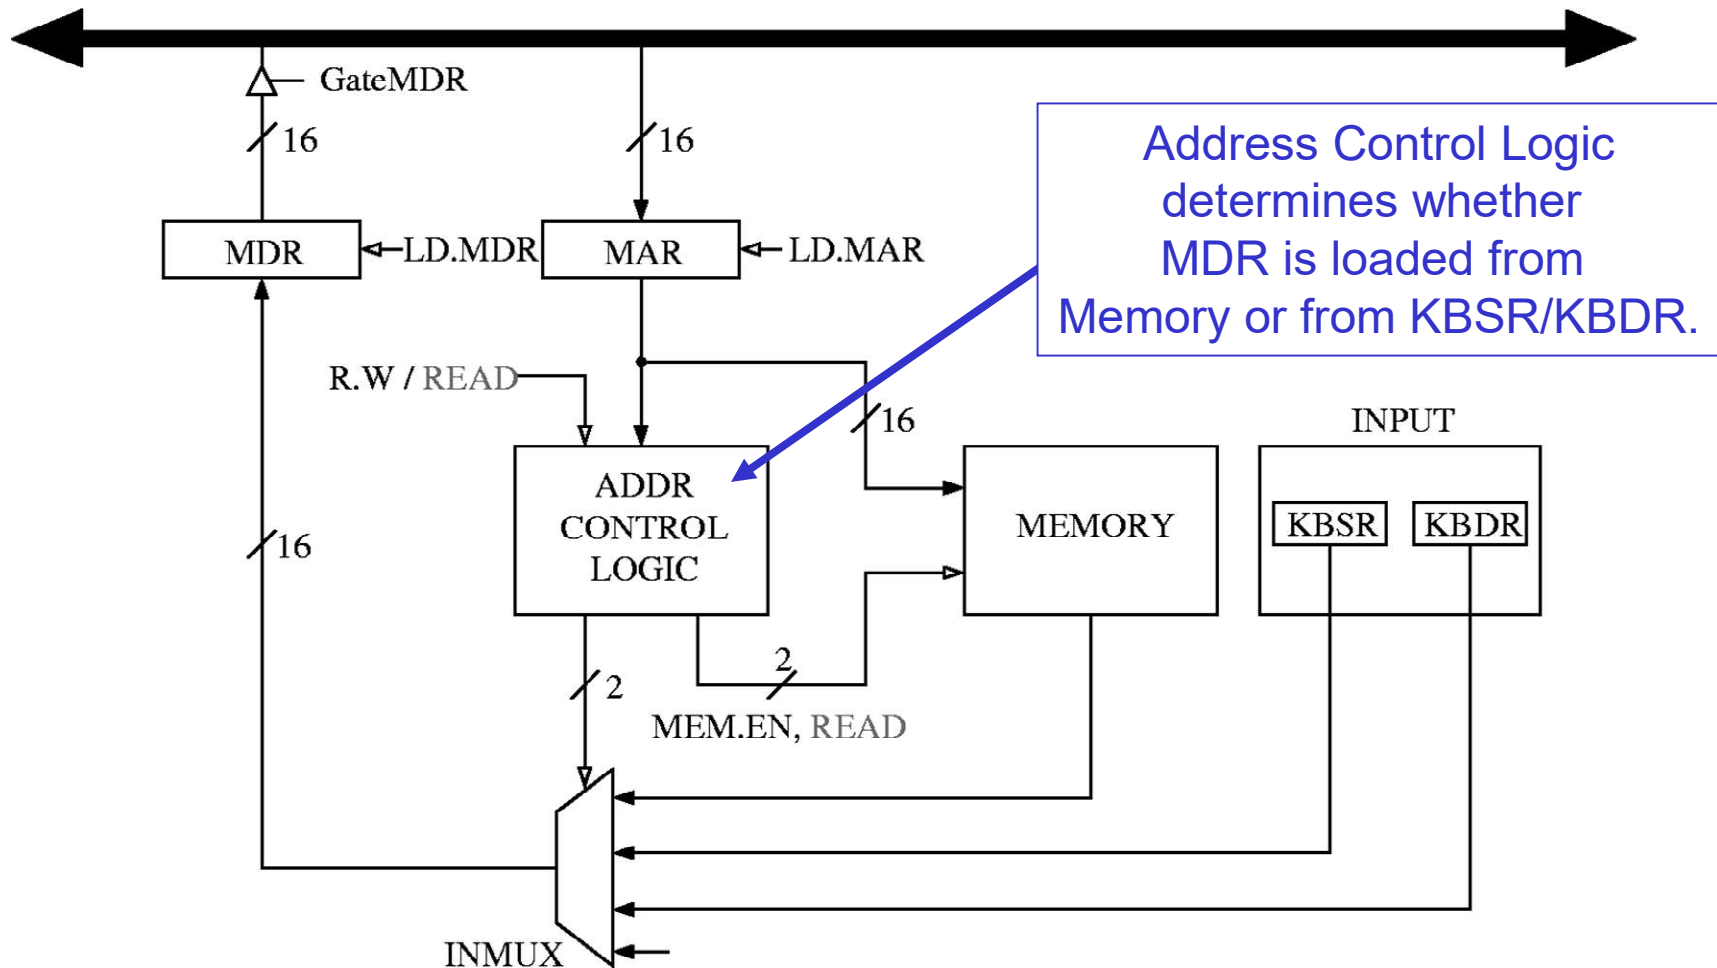
Basic Input Routine



```
POLL      LDI    R0, KBSRPtr
           BRzp   POLL
           LDI    R0, KBDRPtr
           ...

KBSRPtr    .FILL  xFE00
KBDRPtr    .FILL  xFE02
```


Simple Implementation: Memory-Mapped Input



Output to Monitor

When Monitor is ready to display another character:


- the “ready bit” (DSR[15]) is set to one



When data is written to Display Data Register:

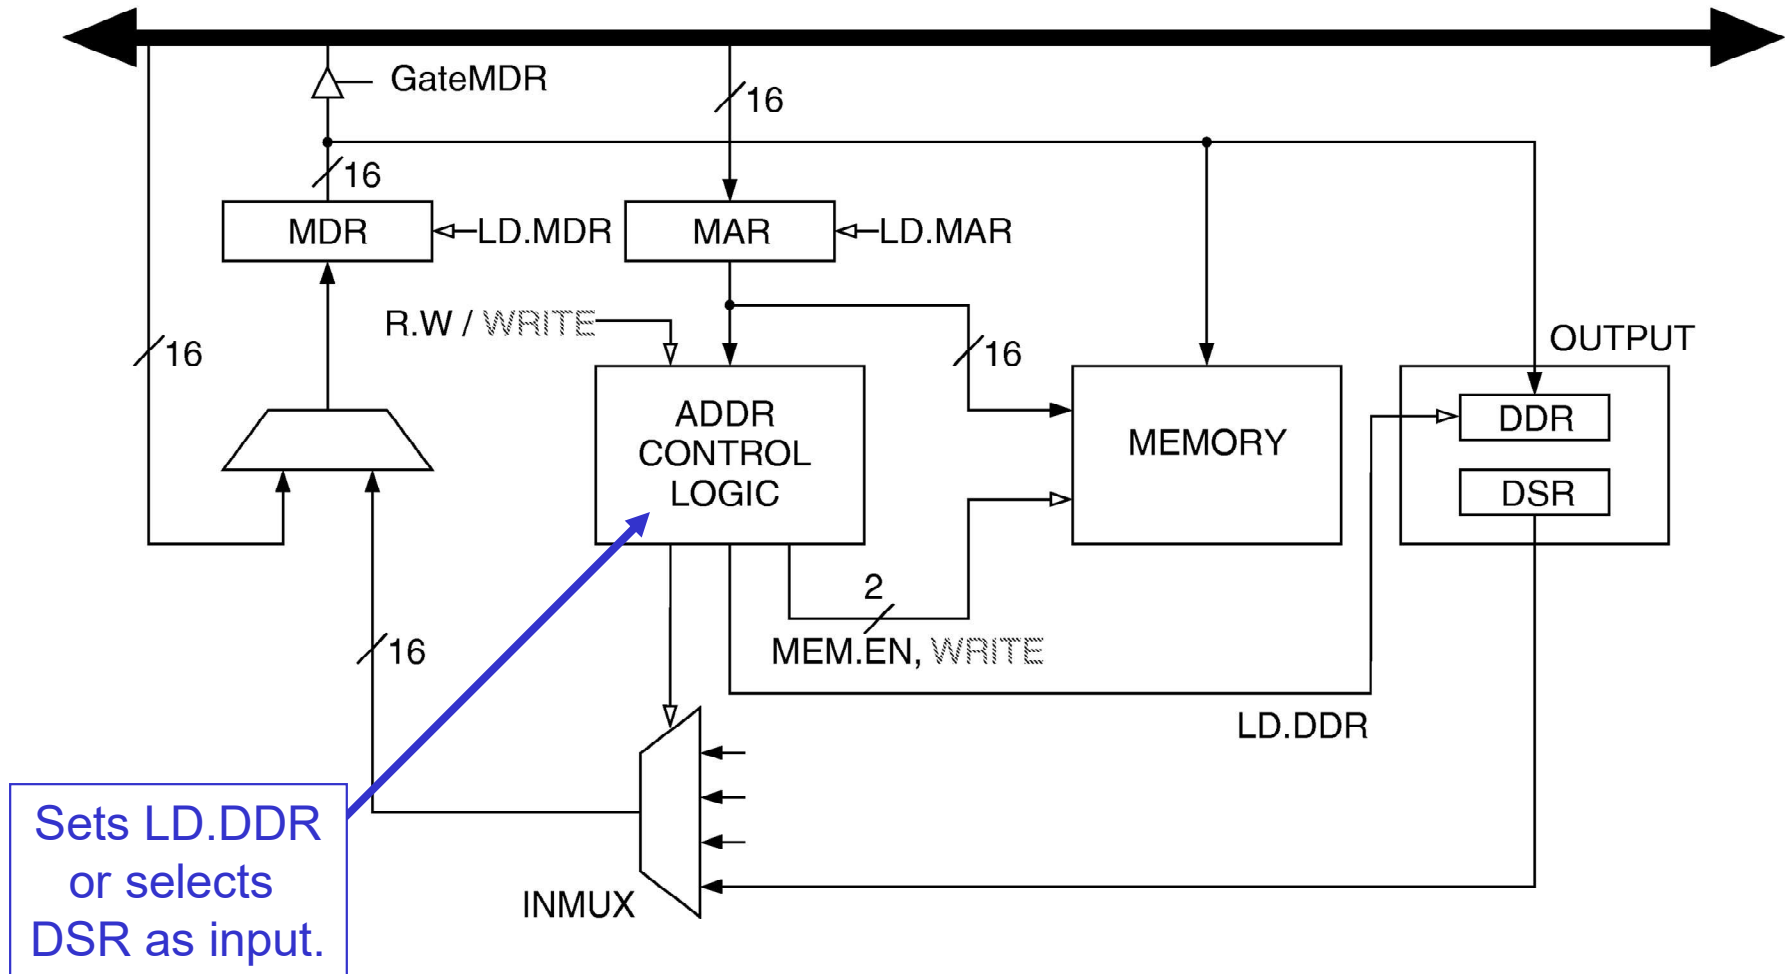
- DSR[15] is set to zero
- character in DDR[7:0] is displayed
- any other character data written to DDR is ignored (while DSR[15] is zero)

Basic Output Routine



```
POLL    LDI    R1, DSRPtr  
        BRzp   POLL  
        STI    R0, DDRPtr  
  
        . . .  
  
DSRPtr  .FILL  xFE04  
DDRPtr  .FILL  xFE06
```

Simple Implementation: Memory-Mapped Output

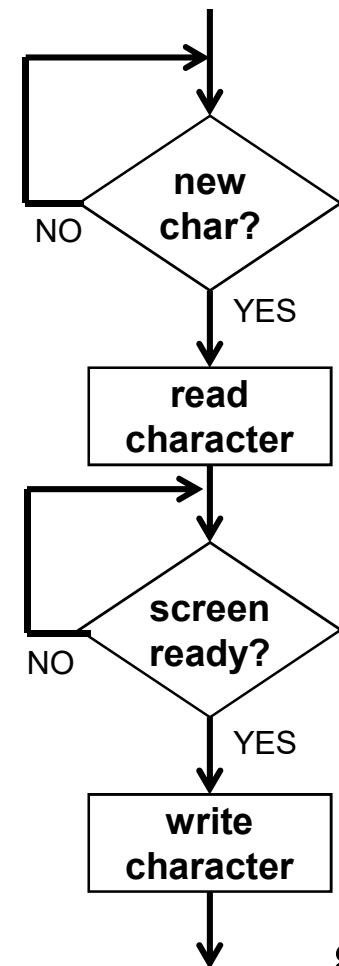


Keyboard Echo Routine

Usually, input character is also printed to screen.

- User gets feedback on character typed and knows its ok to type the next character.

```
POLL1    LDI    R0, KBSRPtr  
          BRzp  POLL1  
          LDI    R0, KBDRPtr  
POLL2    LDI    R1, DSRPtr  
          BRzp  POLL2  
          STI    R0, DDRPtr  
  
          . . .  
  
KBSRPtr  .FILL  xFE00  
KBDRPtr  .FILL  xFE02  
DSRPtr   .FILL  xFE04  
DDRPtr   .FILL  xFE06
```



Polling Overhead

Options for polling

- Don't always need to sit in tight polling loop, doing nothing else.
- But must poll device often enough to not miss any potential data transfers.

Example: mouse

- Must poll mouse at least 30 times per second, to see if (a) change of position and/or (b) button pressed/released.
- How much overhead does this represent?
 - In other words, what percentage of CPU cycles will be used up checking for mouse events?

Mouse Overhead -- Polling

Assume the following parameters:

- **CPU clock is 1 GHz (1×10^9 cycles/sec)**
- **time to service device is 400 cycles, including:**
 - transfer to polling routine, checking status, reading data, updating cursor, returning to user code

How much overhead?

- **time spent polling = $30 \times 400 = 12,000$ cycles/sec**
- **fraction of CPU cycles = $12 \times 10^3 / 1 \times 10^9 = 0.00012\%$**
- **Conclusion: Polling is not a lot of overhead in this case.**
- Not counted -- how to do we know when to poll? Timer interrupt?

Network Overhead -- Polling

What about a network device?

- 100 Mb/s Ethernet, each "packet" is 40B
- Have to poll often enough to not drop data.
- 1 GHz processor, 400 cycles to poll + read data

How much overhead?

- $(320 \text{ bits} / 10^8 \text{ bits/sec}) \times (10^9 \text{ cycles/sec})$
= 1 packet every 3200 cycles = 312,500 packets/sec
- fraction of CPU cycles
= $(400 * 312,500 \text{ cycles/sec}) / (10^9 \text{ cycles/sec}) = 12.5\%$
- This seems a little much -- will interrupts help?

Interrupt-Driven I/O

External device can:

- (1) Force currently executing program to stop;**
- (2) Have the processor satisfy the device's needs; and**
- (3) Resume the stopped program as if nothing happened.**

Why?

- **Polling consumes a lot of cycles, especially for rare events – these cycles can be used for more computation.**
- **Example: Process previous input while collecting current input. (See Example 8.1 in text - page 221.)**

Network Overhead -- Interrupts

Why interrupts?

- Network is not "active" 100% of the time.
Many times, we'll poll and find out there's no data to get.
- Interrupt will happen ONLY when there's actually a packet to be transferred.
- Let's assume network is active 10% of the time.

Overhead?

- $312,500 \text{ packets/sec} \times 10\% = 31,250 \text{ packets/sec}$
- fraction of CPU cycles
 $= (400 \times 31,250 \text{ cycles/sec}) / (10^9 \text{ cycles/sec}) = 1.25\%$

Priority

Every instruction executes at a stated level of urgency.

LC-3: 8 priority levels (PL0-PL7)

- **Example:**
 - **Payroll program runs at PL0.**
 - **Nuclear power correction program runs at PL6.**
- **It's OK for PL6 device to interrupt PL0 program, but not the other way around.**

Priority encoder selects highest-priority device, compares to current processor priority level, and generates interrupt signal if appropriate.


Interrupt-Driven I/O

To implement an interrupt mechanism, we need:

- A way for the I/O device to **signal** the CPU that an interesting event has occurred.
- A way for the CPU to **test** whether the **interrupt signal is set** and whether its **priority is higher** than the current program.

Generating Signal

- Software sets "interrupt enable" bit in device register.
- When ready bit is set and IE bit is set, interrupt is signaled.

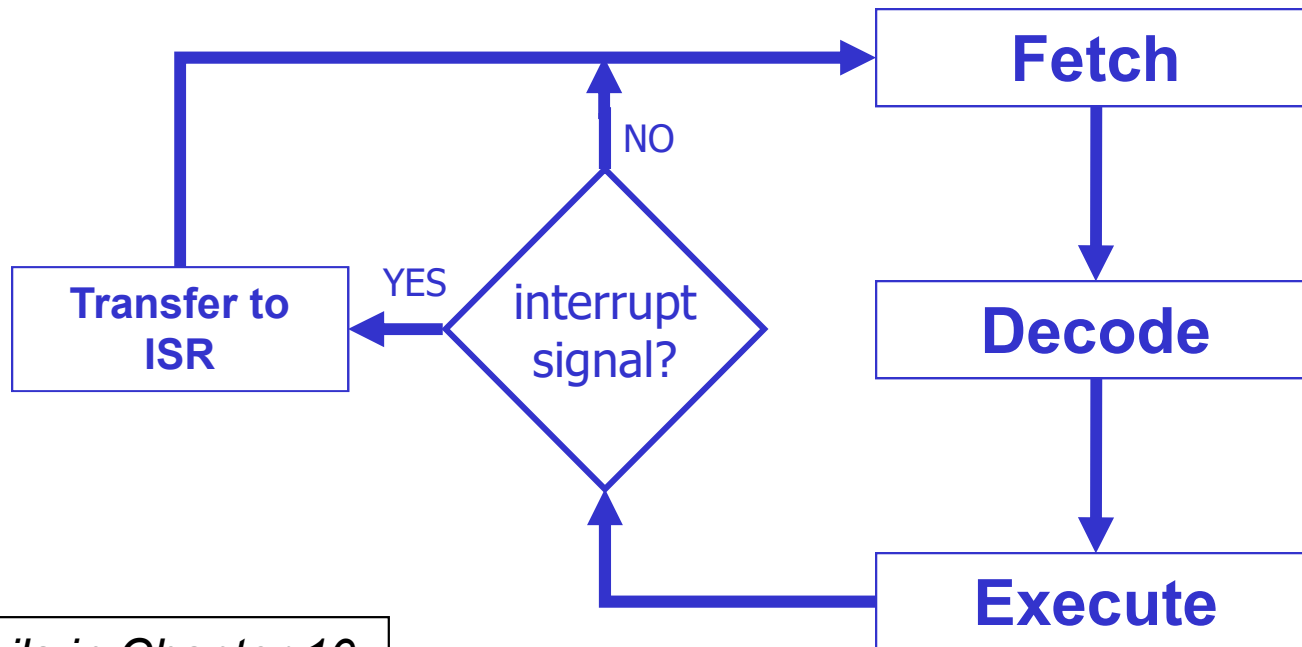


Testing for Interrupt Signal

CPU looks at signal between EXECUTE and FETCH phases.

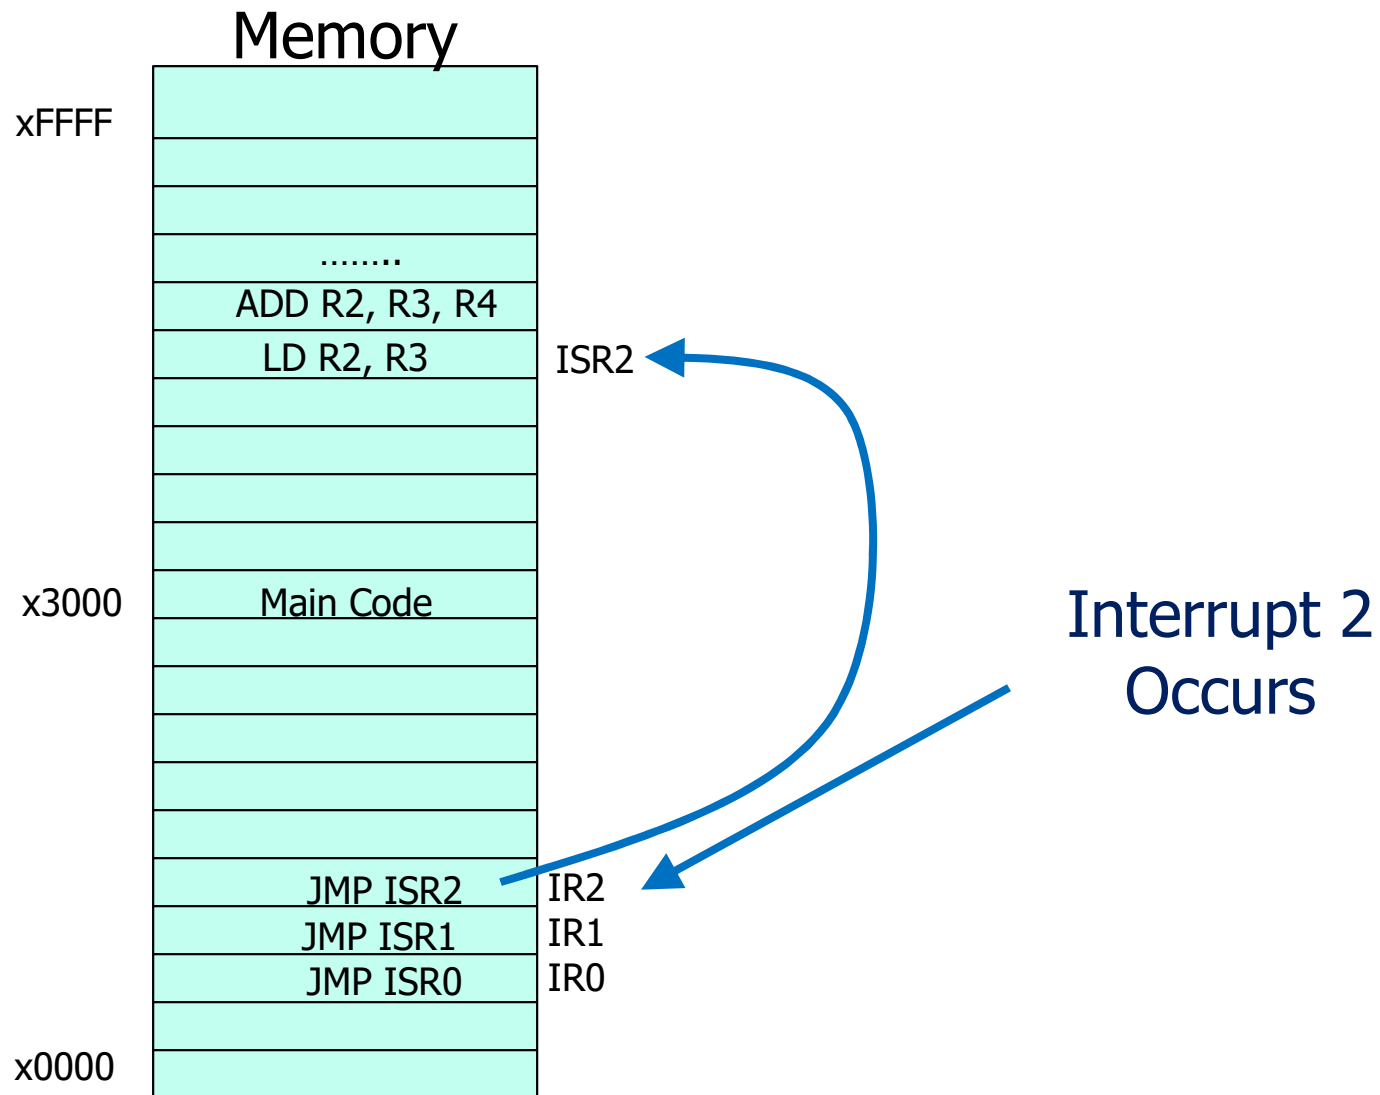
If not set, continues with next instruction.

If set, transfers control to interrupt service routine.

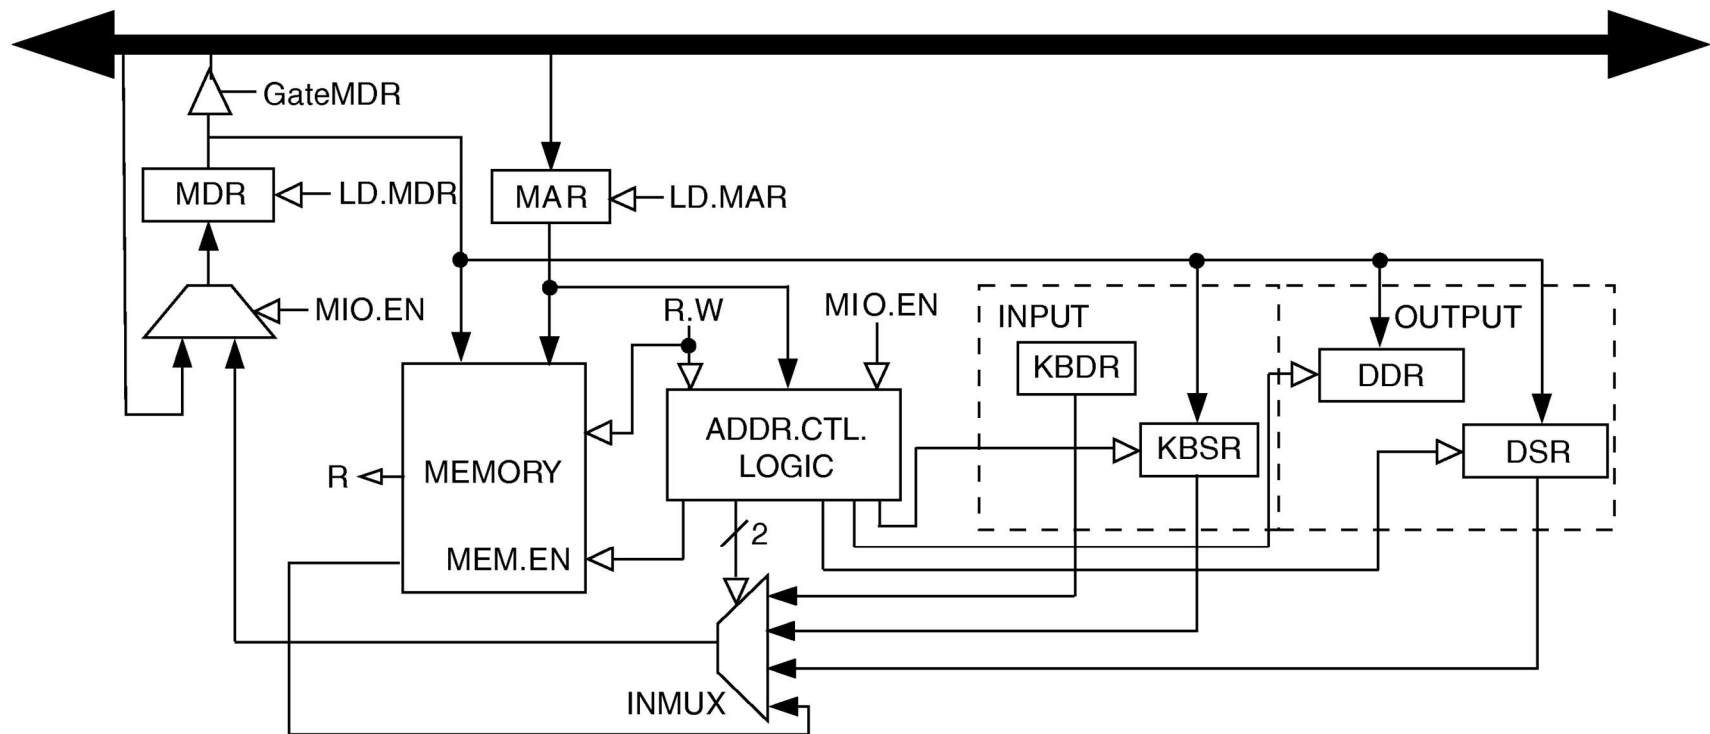


More details in Chapter 10.

Interrupt Service Routines (ISR)



Full Implementation of LC-3 Memory-Mapped I/O



Because of interrupt enable bits, status registers (KB SR/DSR) must be written, as well as read.